

Announcements



If you missed this lecture and are reading these slides online, you need to understand that these are not meant as a stand-alone resource but as an accompaniment to attending the class. If you need to learn the material because you missed the class, do the readings in the Kerrisk text.

How do we communicate between processes?




In order to use a pipe for interprocess communication what has to be true?

1. The process writing to the pipe must be the parent of the process reading from the pipe.
2. The process writing to the pipe must be the child of the process reading from the pipe.
3. The processes must be a parent/child combination but the direction of the pipe could go either way.
4. The processes must be on the same machine.

In order to use a pipe for interprocess communication what has to be true?

1. The process writing to the pipe must be the parent of the process reading from the pipe.
2. The process writing to the pipe must be the child of the process reading from the pipe.
3. The processes must be a parent/child combination but the direction of the pipe could go either way.
4. **The processes must be on the same machine.**

Sockets



- Once you set up the socket, the interface behaves like a file (or pipe)
 - entry in the file descriptor table
 - read/write/close
- Except — sockets are full-duplex (2 way)

Really? Across the Net?

- network protocols
- layers
- clients & servers
- addresses, ports, TCP/IP
- packets
- routing

Relax! Because of the layers of protocols you don't need to know all the details of everything.

Protocols

- Computers use several layers of general protocols to communicate.
- To understand why these layers are important, think about how a company sends you an invoice for a purchase.

TCP/IP


- Transmission Control Protocol.
- Tells us how to package up the data.

source address		dest. address
bytes	ack	port
data		

Details

- make packets

01100111001001
00100010001111
10100010111







- Each TCP packet is given a header
 - sequence number
 - checksum

101010001
111010101
100110010
110101111
001011011

101010001
111010101
100110010
110101111
001011011

101010001
111010101
100110010
110101111
001011011

101010001
111010101
100110010
110101111
001011011



- put in an IP envelope with another header

To
24.197.0.67

To
24.197.0.67

To
24.197.0.67

To
24.197.0.67

Client Server Model



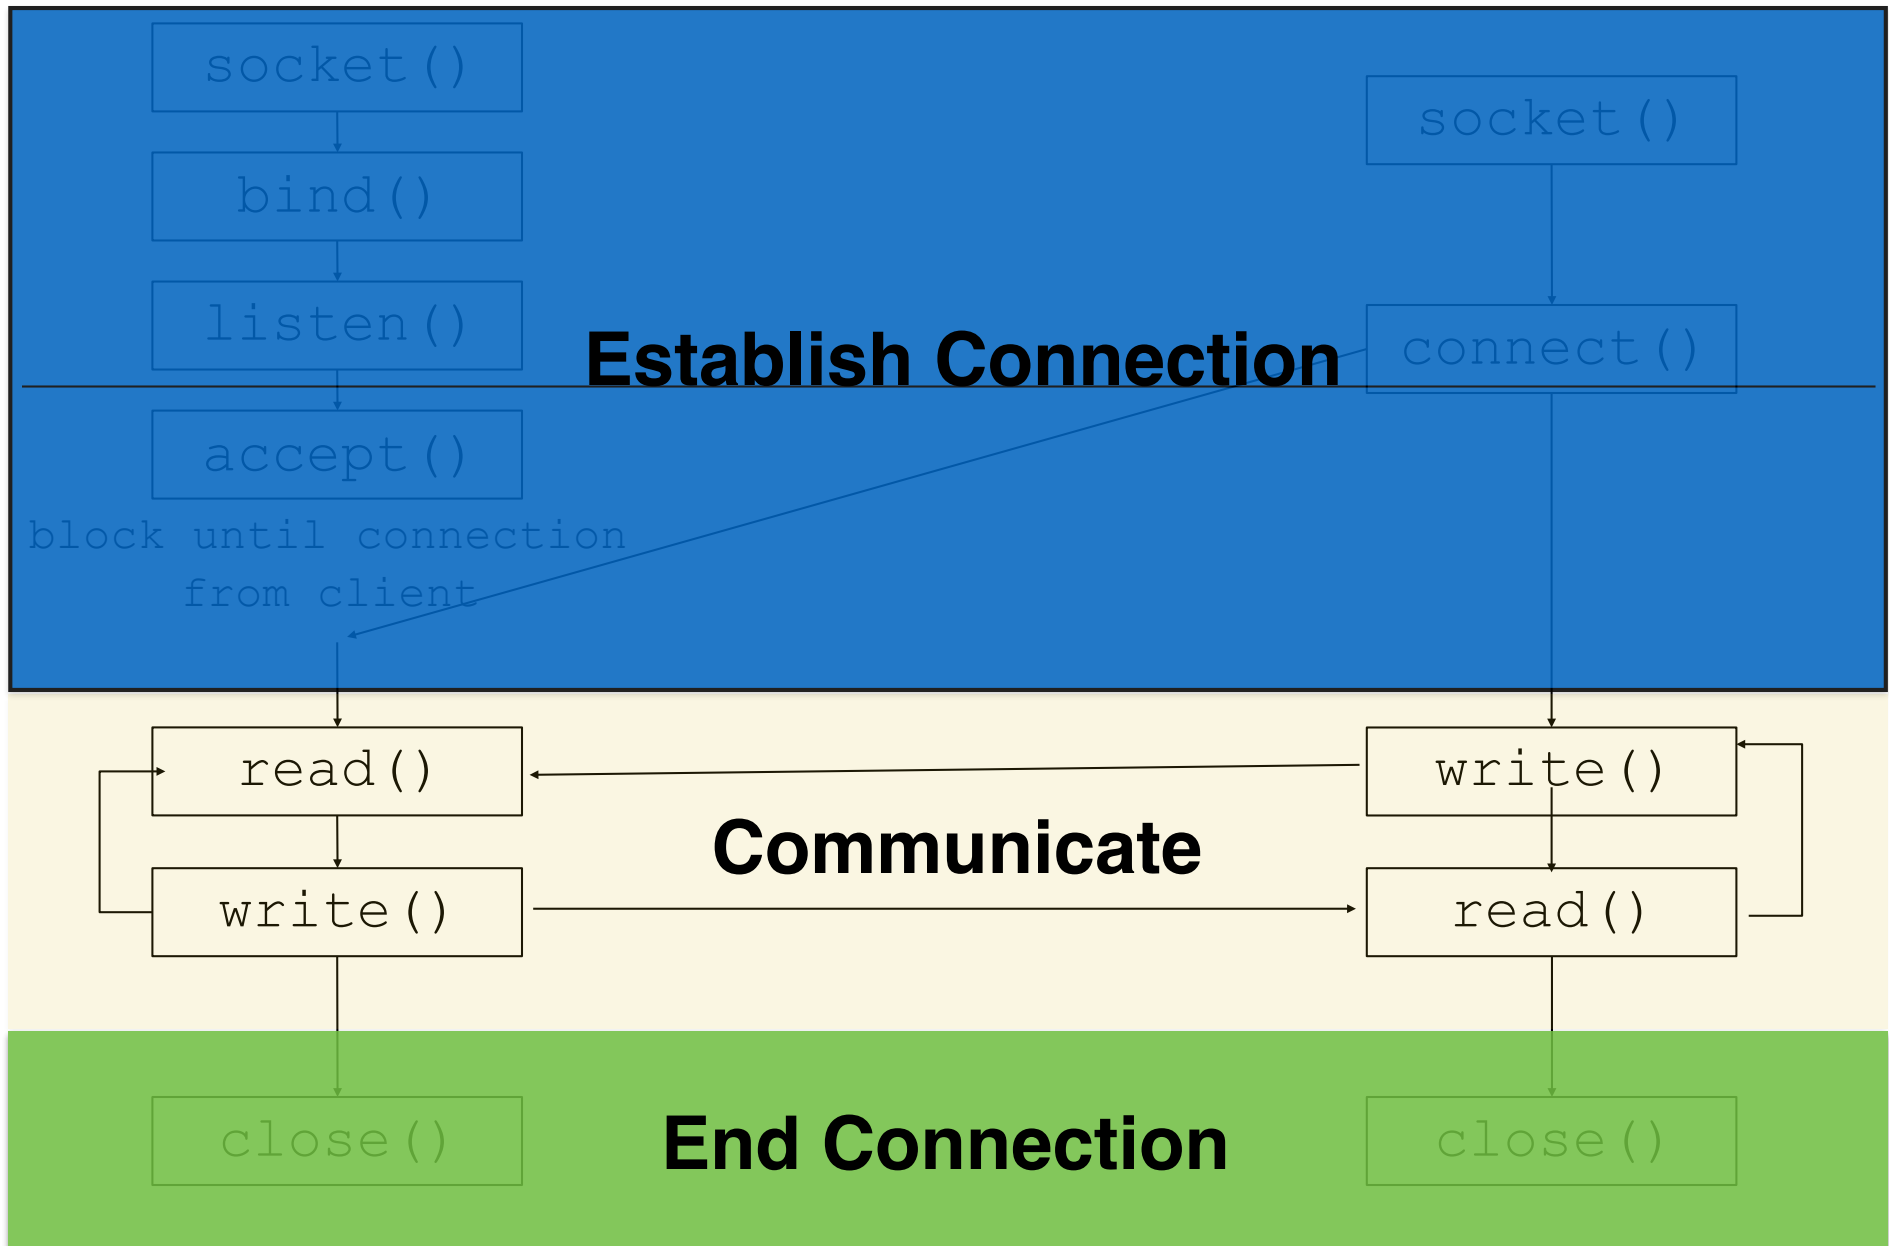
Server: Always available at a known location



Client: Shows up and initiates the connection

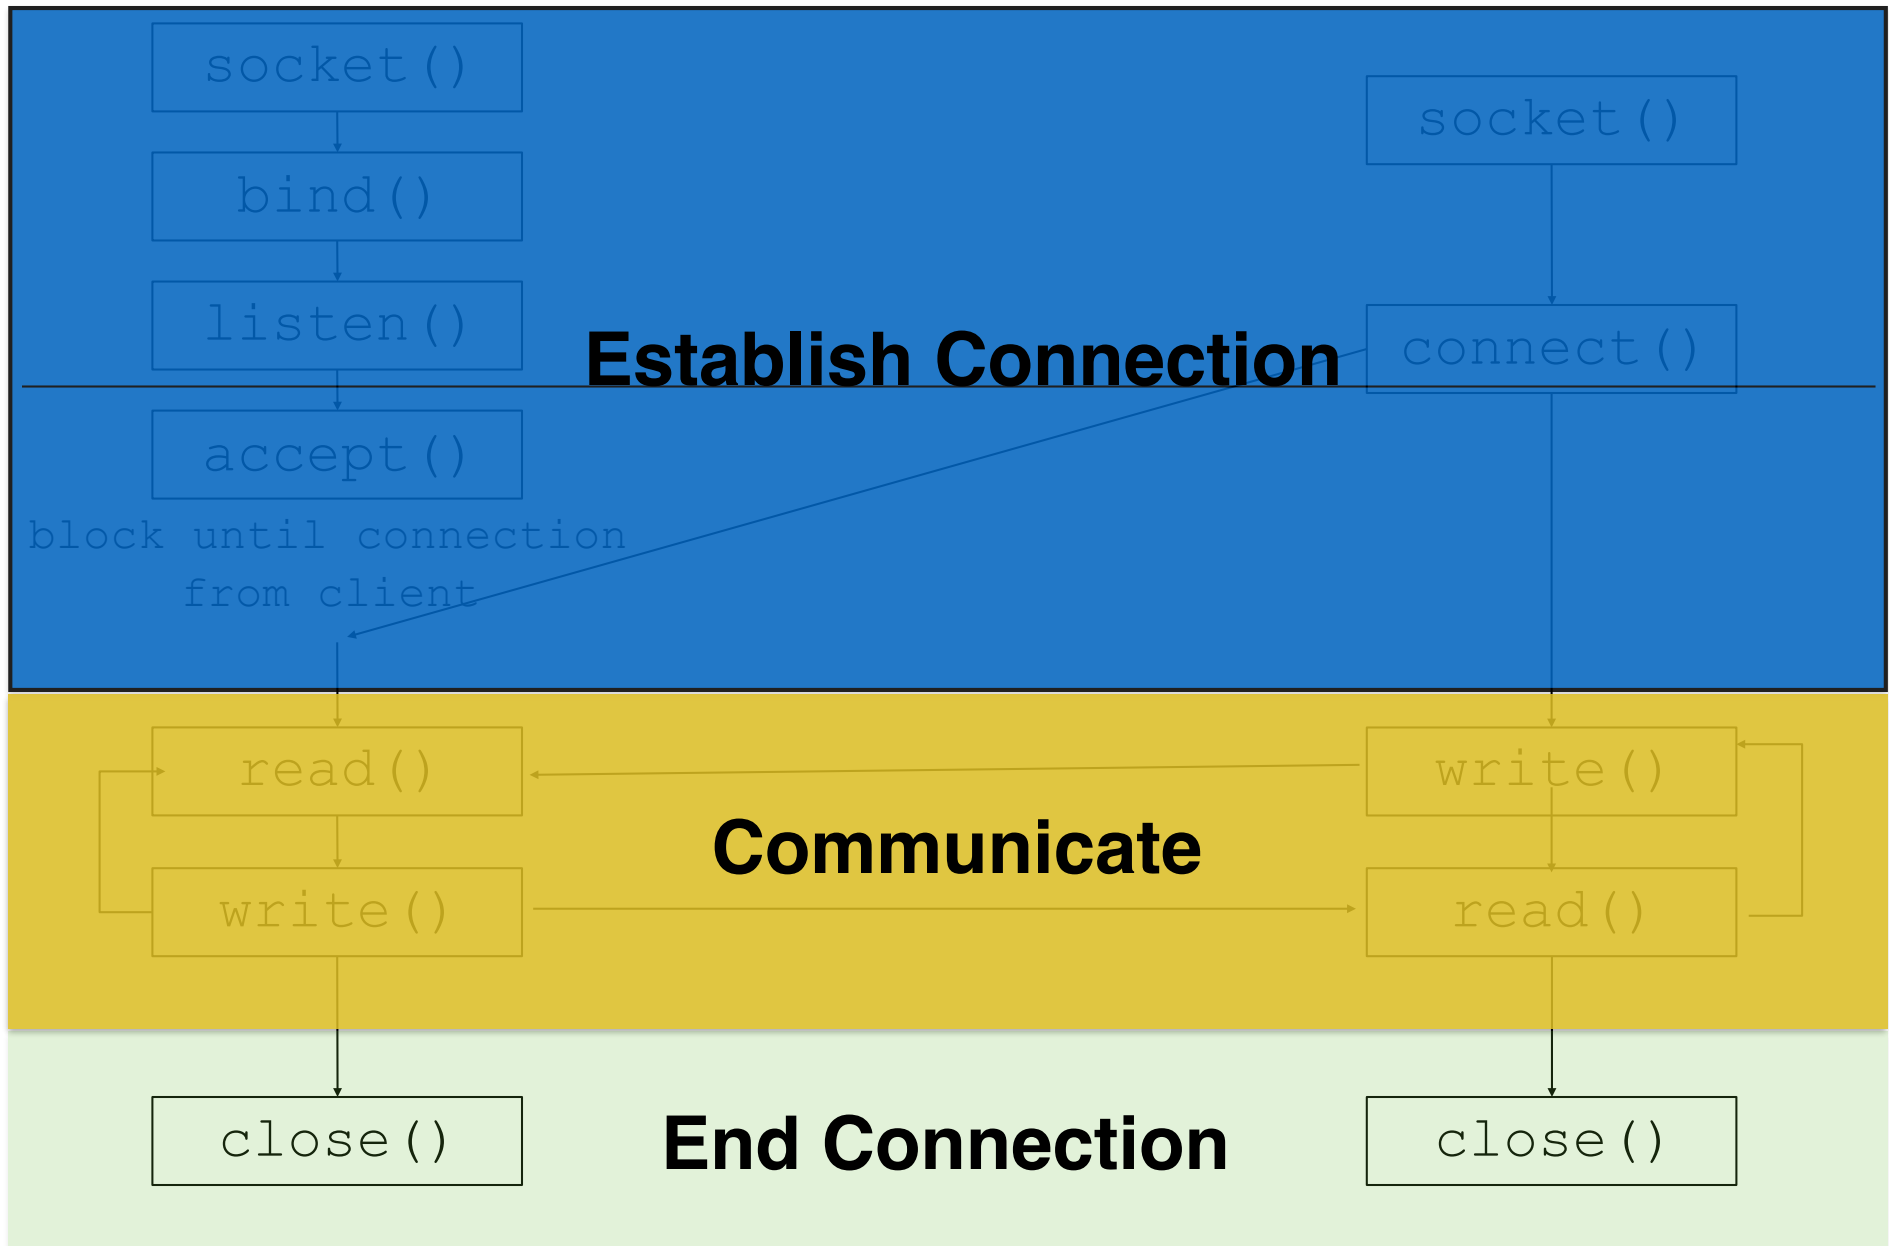
Server

Client



Server

Client



Server

Client

socket()

bind()

设置socket的地址

listen()

set up queue for client access request

accept()

看queue有没有请求，如果有处理，没有等待

block until connection
from client

read()

write()

close()

socket()

connect()

write()

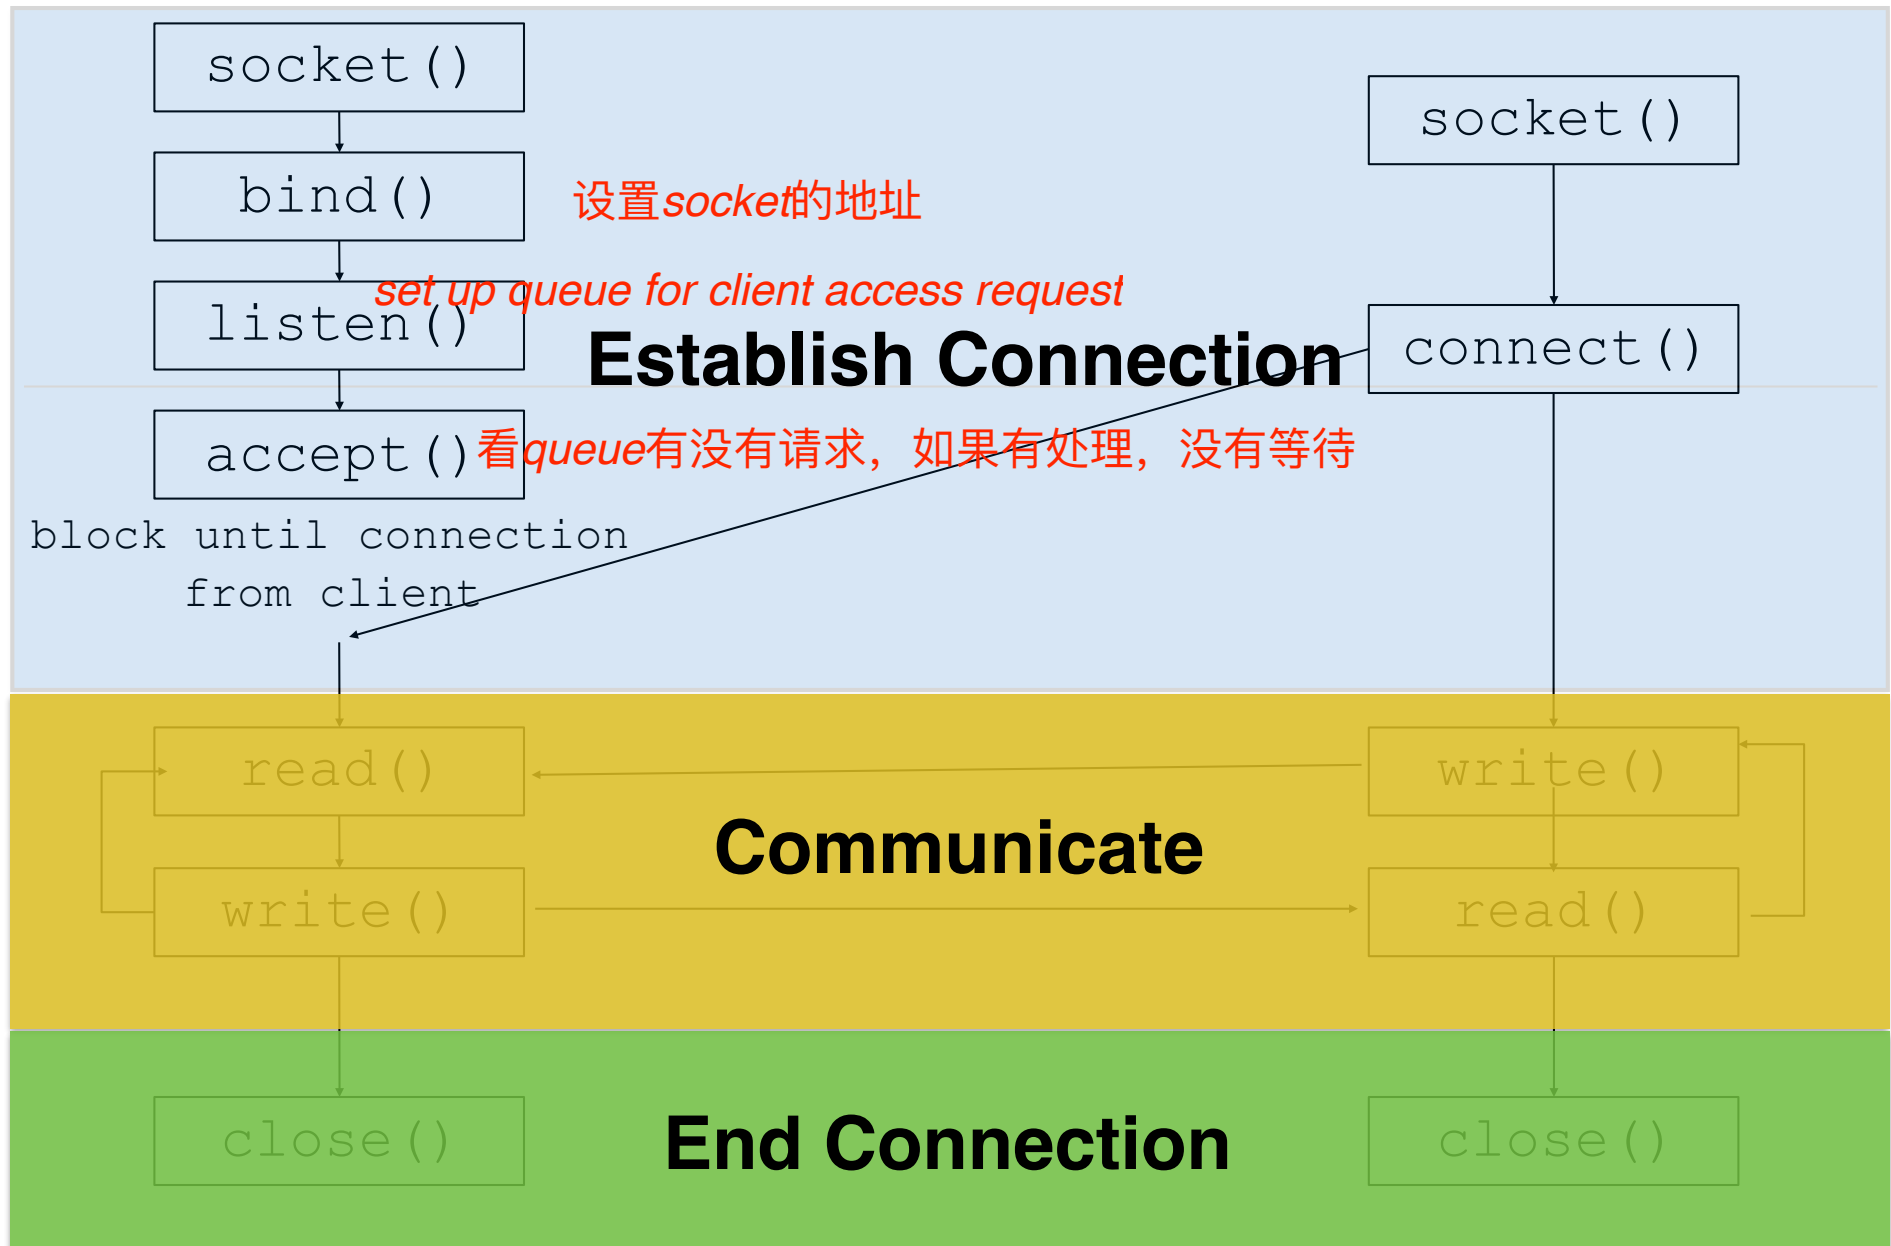
read()

close()

Establish Connection

Communicate

End Connection



Learning More at Home

- We are only talking about connection-oriented sockets in CSC 209
- If you want to know about other sockets, read chapters 56-61 in Kerrisk
- If you missed this class and are trying to simply read these slides, read the chapters in Kerrisk listed on the course website

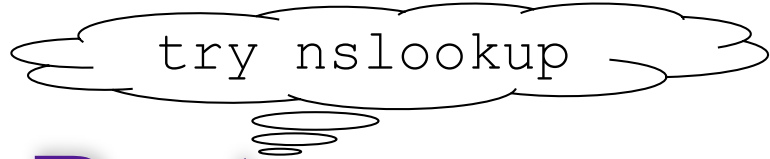
Connection-Oriented

Server

- Create a socket: `socket()`
- Assign an address to a socket: `bind()`
- Establish a queue for connections: `listen()`
- Get a connection from the queue: `accept()`

Client

- Create a socket: `socket()`
- Initiate a connection: `connect()`



try nslookup

Addresses and Ports

- A **socket pair** is the two endpoints of the connection.
- An endpoint is identified by an **IP address and a port**.
- IPv4 addresses are 4 8-bit numbers:
 - 128.100.31.200 = wolf.teach.cs.toronto.edu
- Ports
 - because multiple processes can communicate with a single machine we need another identifier.

More on Ports

- Well-known ports: 0-1023
 - 80 = http
 - 21 = ftp
 - 22 = ssh
 - 25 = smtp (mail)
 - 23 = telnet
 - 194 = irc
- Registered ports: 1024-49151
 - 2709 = supermon
 - 26000 = quake
 - 3724 = world of warcraft
- Dynamic (private) ports: 49152-65535
 - You should pick ports in this range to avoid overlap

Server

Client

socket()

bind()

listen()

accept()

block until connection
from client

read()

write()

close()

socket()

connect()

write()

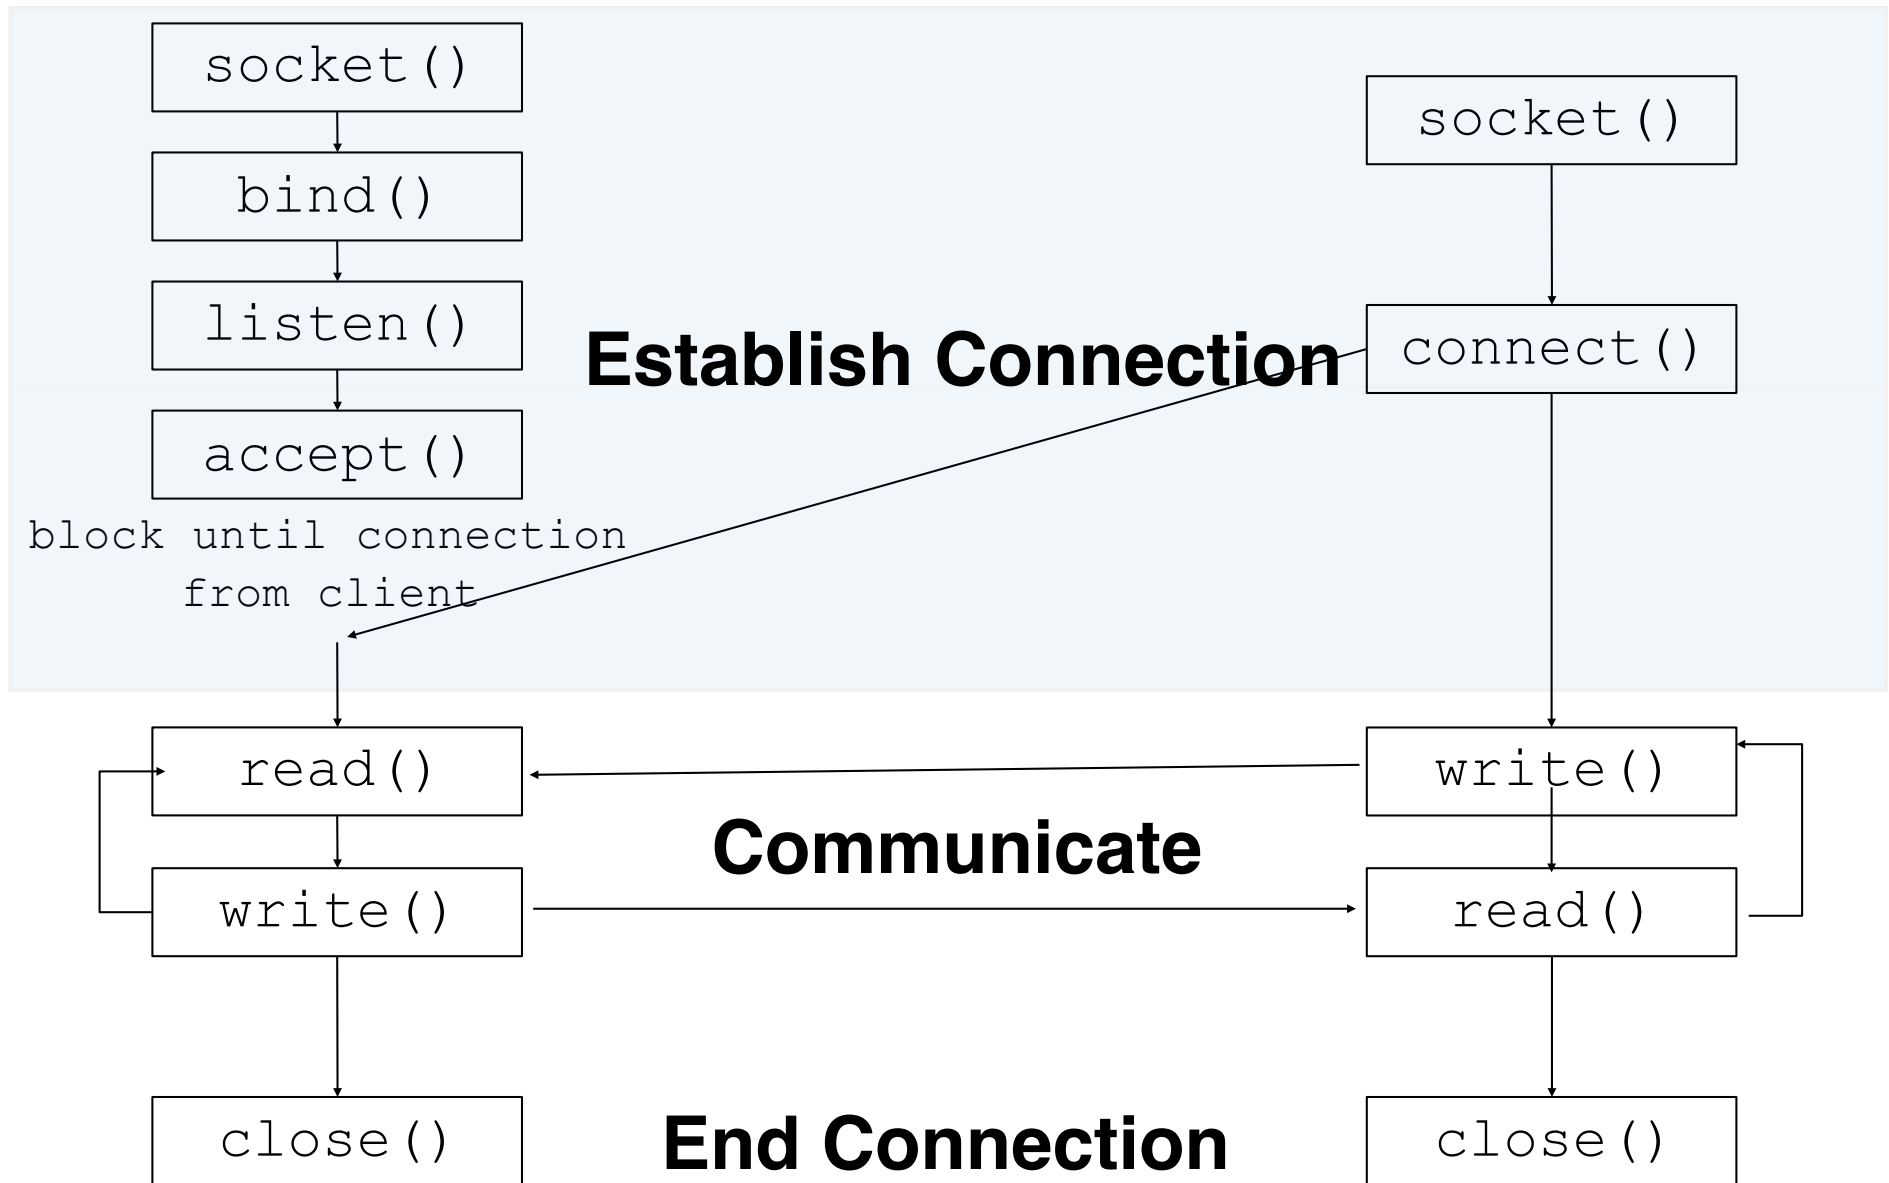
read()

close()

Establish Connection

Communicate

End Connection



Server side

```
int socket(int family, int type,  
           int protocol);
```

- family specifies protocol family:
 - `AF_INET` – IPv4
 - `AF_LOCAL` – Unix domain
- type
 - `SOCK_STREAM`, `SOCK_DGRAM`, `SOCK_RAW`
- protocol
 - set to `0` except for RAW sockets
- returns a socket descriptor

bind to a name

```
int bind(int sockfd,  
         const struct sockaddr *servaddr,  
         socklen_t addrlen);
```

- **sockfd** – returned by **socket**
- **struct sockaddr_in** {
 short sin_family; /*AF_INET */
 u_short sin_port;
 struct in_addr sin_addr;
 char sin_zero[8]; /*filling*/
};
- **sin_addr** can be set to **INADDR_ANY** to communicate on any network interface.

Set up queue in kernel

```
int listen(int sockfd, int backlog)
```

- after calling `listen`, a socket is ready to accept connections
- prepares a queue in the kernel where partially completed connections wait to be accepted.
- `backlog` is the maximum number of partially completed connections that the kernel should queue.

Complete the connection

```
int accept(int sockfd,  
          struct sockaddr *cliaddr,  
          socklen_t *addrlen);
```

- blocks waiting for a connection (from the queue)
- returns a new descriptor which refers to the TCP connection with the client
- `sockfd` is the listening socket
- `cliaddr` is the address of the client
- reads and writes on the connection will use the socket returned by `accept`

Client side

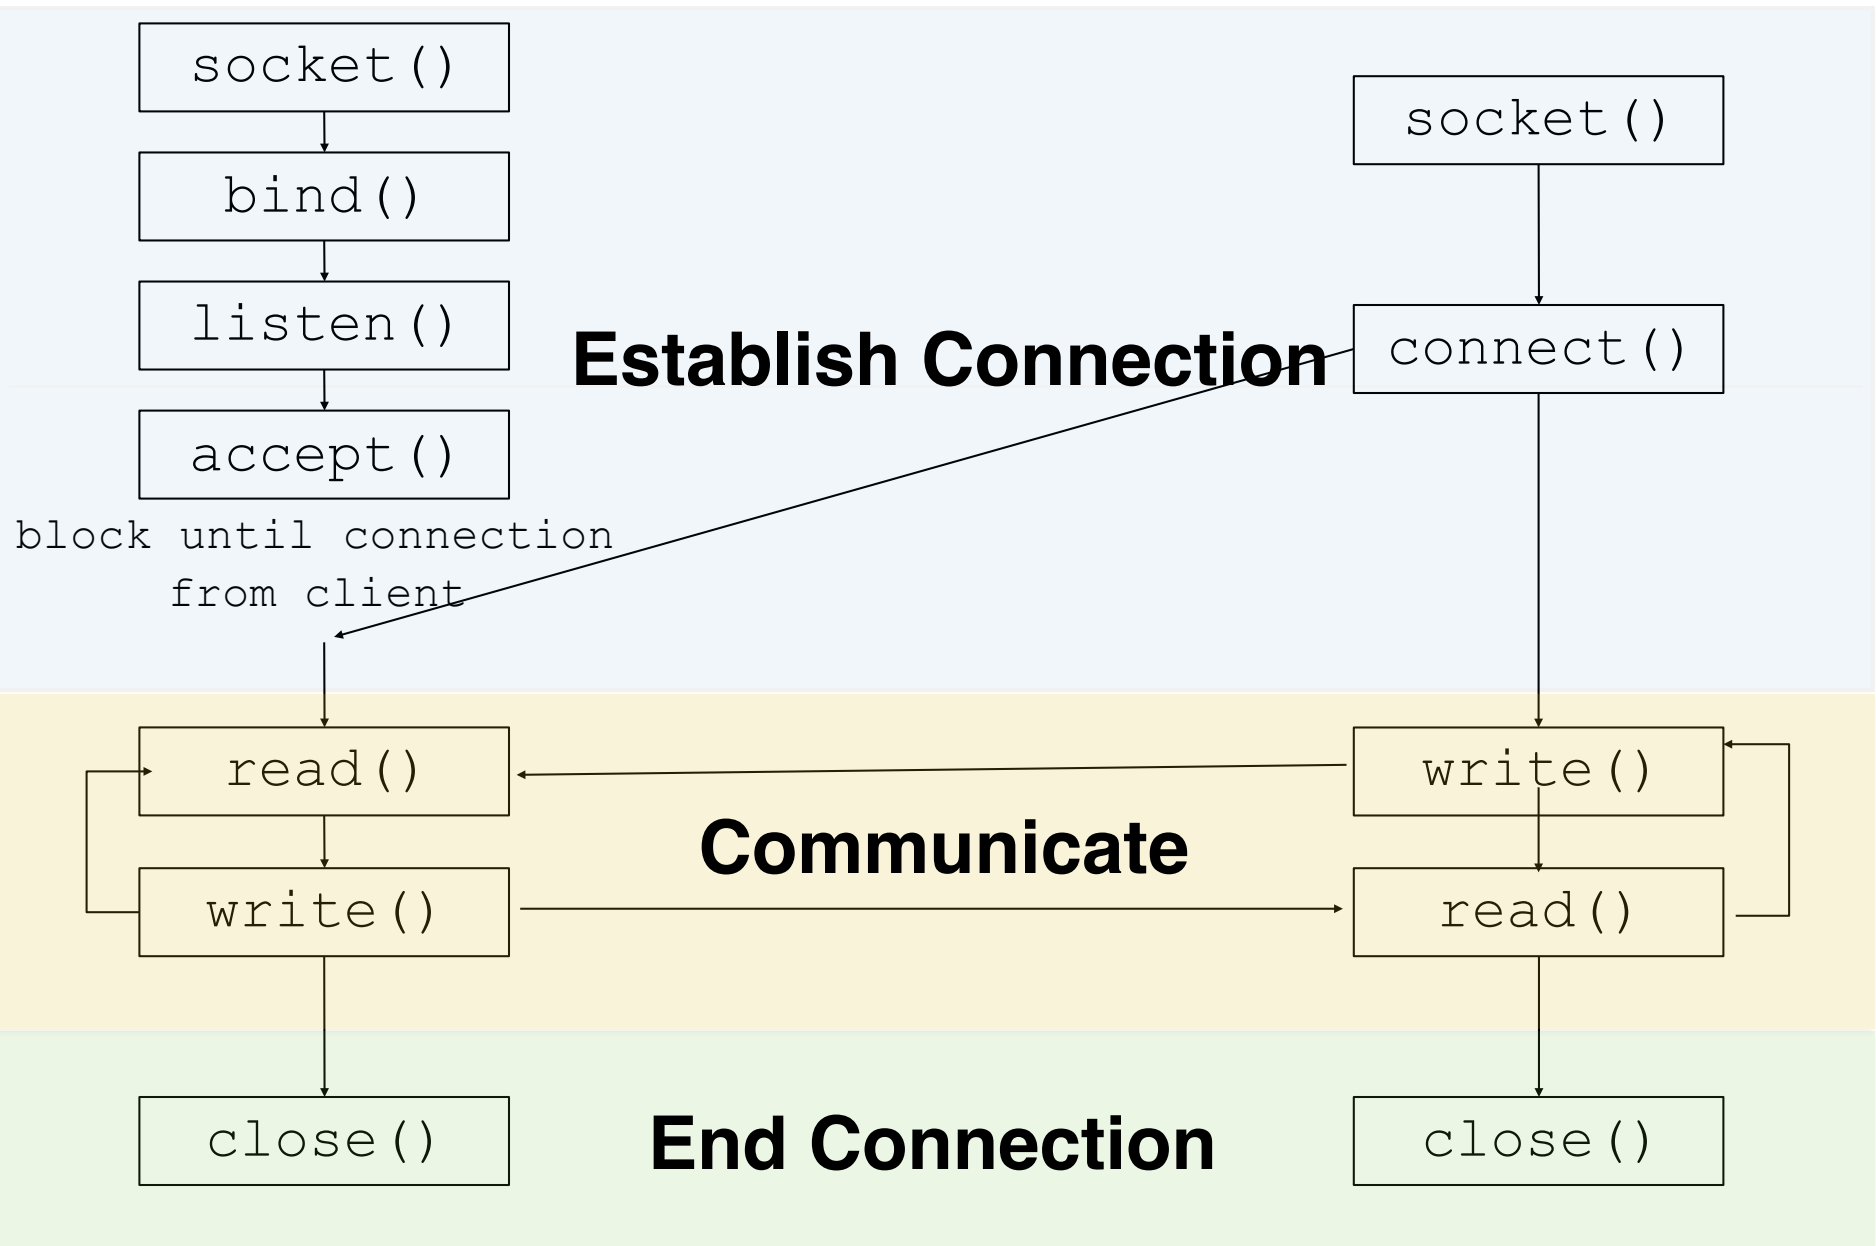
- `socket()` – same as server, to say “how” we are going to talk

```
int connect(int sockfd,  
            const struct sockaddr *servaddr,  
            socklen_t addrlen);
```

- the kernel will choose a dynamic port and source IP address.
- returns 0 on success and -1 on failure setting `errno`.


Server

Client




Byte order

- Big-endian



- Little-endian



- Intel is little-endian, and Sparc is big-endian

Network byte order

- To communicate between machines with unknown or different “endian-ness” we convert numbers to network byte order (big-endian) before we send them.
- There are functions provided to do this:
 - `unsigned long htonl(unsigned long)`
 - `unsigned short htons(unsigned short)`
 - `unsigned long ntohl(unsigned long)`
 - `unsigned short ntohs(unsigned short)`

Network Newline

- `\r\n` rather than just `\n`

Helpful Tips

- Think carefully about the exact bytes you are sending and how the receiver will interpret them
- Pay attention to ends of strings and extra characters in char arrays
- Byte order & network newlines
- Don't assume full lines will arrive in a single read (practice in this week's lab)