



Balanced Trees: AVL Tree

Fatemeh Panahi
Department of Computer Science
University of Toronto
CSC263-Fall 2017
Lecture 4

Announcements

- A1 Solutions are out, A2 handout will be posted on Friday.
- **Tutorial on Friday:**
Textbook exercise for tutorial:
12.2-6, 12.2-7, 12.2-8, 12-2 and problem 6-3.
Work on these exercises before the tutorial.
Some examples on AVL Tree and Augmentation.

Today

- Augmentation
- AVL Tree (Not covered in CLRS)
 - Height and Balance factor
 - Operations
 - Search
 - Insert
 - Delete

Augmented Data Structures:

existing data structure **modified** to store additional information and/or perform additional operations.

Balanced BST

AVL tree, Red-Black tree, 2-3 tree, AA tree,



Reflect on AVL tree

- We “augmented” BST by storing additional information (the balance factor) at each node.
- The additional information enabled us to do additional cool things with the BST (keep the tree balanced).
- And we can maintain this additional information efficiently in modifying operations (within $O(\log n)$ time, without affecting the running time of Insert or Delete).

Augmentation: General Procedure

1. Choose data structure to augment
2. Determine additional information
3. Check additional information can be maintained, during each original operation, hopefully efficiently.
4. Implement new operations.

AVL Trees

Why AVL tree?

- First self-balancing BST to be invented.

Why is it called AVL?


- Invented by Georgy **A**delson-**V**elsky and E. M. **L**andis in 1962.

Height of a BST Node

$x.height \leftarrow 1 + \max(x.left.height, x.right.height)$

Trick: make NIL be an actual node with

- . NIL.item = NIL,
- . NIL.left = NIL,
- . NIL.right = NIL,
- . **NIL.height = -1.**



AVL Trees


An extra attribute to each node in a BST: **balance factor**

$h_R(x)$: height of x 's right subtree

$h_L(x)$: height of x 's left subtree



$$\underline{BF(x) = h_R(x) - h_L(x)}$$

- $BF(x) = 0$: x is balanced
- $BF(x) = 1$: x is right-heavy
- $BF(x) = -1$: x is left-heavy
- $BF(x) > 1$ or < -1 : x is imbalanced





AVL tree: definition

An AVL tree is a BST in which every node is balanced, right-heavy or left-heavy. i.e., the BF of every node must be 0, 1 or -1.



AVL tree: definition

An AVL tree is a BST in which every node is balanced, right-heavy or left-heavy. i.e., the BF of every node must be 0, 1 or -1.



AVL Trees - height

If there are n nodes, what is the maximum possible height?

\Leftrightarrow

If the height is h , what is the minimum possible number of nodes?

Let $M(h)$ be the **minimum number of nodes in any AVL tree** of height h . Then $M(0) = 1, M(1) = 2$, and $M(h) = 1 + M(h - 1) + M(h - 2)$ for $h \geq 2$.

$$fib(0) = 0, fib(1) = 1, fib(2) = 1, fib(3) = 2, \quad fib(h) = \frac{\phi^h - (1-\phi)^h}{\sqrt{5}} \quad \phi = \frac{1+\sqrt{5}}{2} \approx 1.61$$
$$M(h) = fib(h + 2) - 1$$

AVL Trees - height

$$fib(h) = \frac{\phi^h - (1-\phi)^h}{\sqrt{5}} \quad \phi = \frac{1+\sqrt{5}}{2} \approx 1.61, 1 - \phi \approx -0.61$$

$$M(h) = fib(h + 2) - 1$$

$$n \geq M(h) = \frac{(\phi^{h+2} - (1-\phi)^{h+2})}{\sqrt{5}} - 1 > \frac{\phi^{h+2}}{\sqrt{5}} - 1 - 1$$

$$\frac{\phi^{h+2}}{\sqrt{5}} - 2 < n \rightarrow \phi^{h+2} < \sqrt{5}(n + 2) \rightarrow h + 2 < 1.44 \log n + 2 + constant$$

This implies $h \in O(\log n)$.

Height of an AVL tree with n nodes is $O(\log n)$.

Operations on AVL trees

- AVL-Search(root, k)
- AVL-Insert(root, x)
- AVL-Delete(root, x)

What should we consider for these operations:

- Before the operation, the BST is a valid AVL tree (**precondition**)
- After the operation, the BST must still be a valid AVL tree: so re-balancing may be needed.
- The **balance factor** attributes of some nodes need to be updated.

AVL-Search(root, k)

Search for key k in the AVL tree rooted at root


- First, do a `TreeSearch(root, k)` as in BST.
- Then, nothing else!

(The tree is still balanced since we didn't change the tree)

AVL-Insert(root, x)

First, do a TreeInsert(root, x) as in BST, then update the balance factors

Example 1: Insert(12)




Still an AVL Tree,
Nothing has to change!

AVL-Insert(root, x)


First, do a TreeInsert(root, x) as in BST, then update the balance factors

Example 1: Insert(4)



Not an AVL Tree anymore,
Needs rebalancing!

Basic move for rebalancing - Rotation



- Inorder walk of the tree is the same: $A y B x C$
- Three references to change: parent's left/right, y .left, x .right.
- Decreases depth of subtree A , increases depth of subtree C .
- Height is just updated accordingly as rotations happen and nobody outside the picture needs to be updated, because the height is the same as before and nobody above would notice a difference. So, only need $O(1)$ time for updating heights.

Run time:

$\Theta(1)$

Rotation to the Left

AVL-ROTATE-TO-THE-LEFT(x):

Rearrange references.

$y \leftarrow x.\text{right}$

$x.p \leftarrow y$

$x.\text{right} \leftarrow y.\text{left}$

$y.\text{left} \leftarrow x$

Update heights.


AVL-UPDATE-HEIGHT(x)

AVL-UPDATE-HEIGHT(y)

Return new parent.

return y

Left rotation (x)



AVL-UPDATE-HEIGHT(x):

$x.\text{height} \leftarrow 1 + \max(x.\text{left}.\text{height}, x.\text{right}.\text{height})$

Rotation to the right

AVL-ROTATE-TO-THE-RIGHT(y):

Rearrange references.

$x \leftarrow y.\text{left}$

$x \leftarrow y.p$

$y.\text{left} \leftarrow x.\text{right}$

$x.\text{right} \leftarrow y$

$\equiv \uparrow$


Update heights.

AVL-UPDATE-HEIGHT(y)

AVL-UPDATE-HEIGHT(x)

Return new parent.

return x




AVL-UPDATE-HEIGHT(x):

$x.\text{height} \leftarrow 1 + \max(x.\text{left}.\text{height}, x.\text{right}.\text{height})$


When do we need to rebalance?

Let x be the lowest ancestor of the new node who became imbalanced.

Case 1: the insertion increases the height of a node's **right subtree**, and that node was already **right heavy**.





Case 2: the insertion increases the height of a node's **left subtree**, and that node was already **left heavy**.



Case 1

In order to rebalance, we need to increase the height of the left subtree and decrease the height of the right subtree, so we perform a left rotation



case 1.2

case 1.1


The height of subtrees A and B are both h .

Is it possible that height of one is h and height of the other is $h - 1$?

x is the lowest node that become imbalanced.


Case 1.1

Perform a left rotation



case 1.1

Left rotation




The tree becomes balanced after a single rotation.

Remark: After the rotation, the height of the whole subtree in the picture does not change, everything happens in this picture stays in this picture, nobody above would notice.


Case 1.2

Perform a left rotation



case 1.2

Left rotation (x)




Still unbalanced!




Case 1.2.1: more refined graph

Perform a left rotation



Case 1.2.2: more refined graph

Perform a left rotation



AVL-rebalancing

○ Case 1:


- Case 1.1: single **left** rotation
- Case 1.2: double **right-left** rotation

AVL-REBALANCE-TO-THE-LEFT(root)

if *root.right.left.height* > *root.right.right.height*

root.right \leftarrow *AVL-ROTATE-TO-THE-RIGHT(root.right)*

return *AVL-ROTATE-TO-THE-LEFT(root)*



○ Case 2: (symmetric to Case 1)


- Case 2.1: single **right** rotation
- Case 2.2: double **left-right** rotation

AVL-REBALANCE-TO-THE-RIGHT(root)

if *root.left.right.height* > *root.left.left.height*

root.left \leftarrow *AVL-ROTATE-TO-THE-LEFT(root.left)*

return *AVL-ROTATE-TO-THE-RIGHT(root)*



Some examples:

Insert 3, 2, 1, 4, 5, 6, 7, 16, 15, 14



Fig 1




Fig 2




Fig 3

Right rotation
→




Fig 4




Fig 5




Fig 6

Left rotation
→







Fig 7

Example - continue

Insert: 3, 2, 1, 4, 5, 6, 7, 16, 15, 14






Left rotation(2)



Example - continue

Insert: 3, 2, 1, 4, 5, 6, 7, 16, 15, 14



Example - continue

Insert: 3, 2, 1, 4, 5, 6, 7, 16, 15, 14




Fig 16




Fig 17


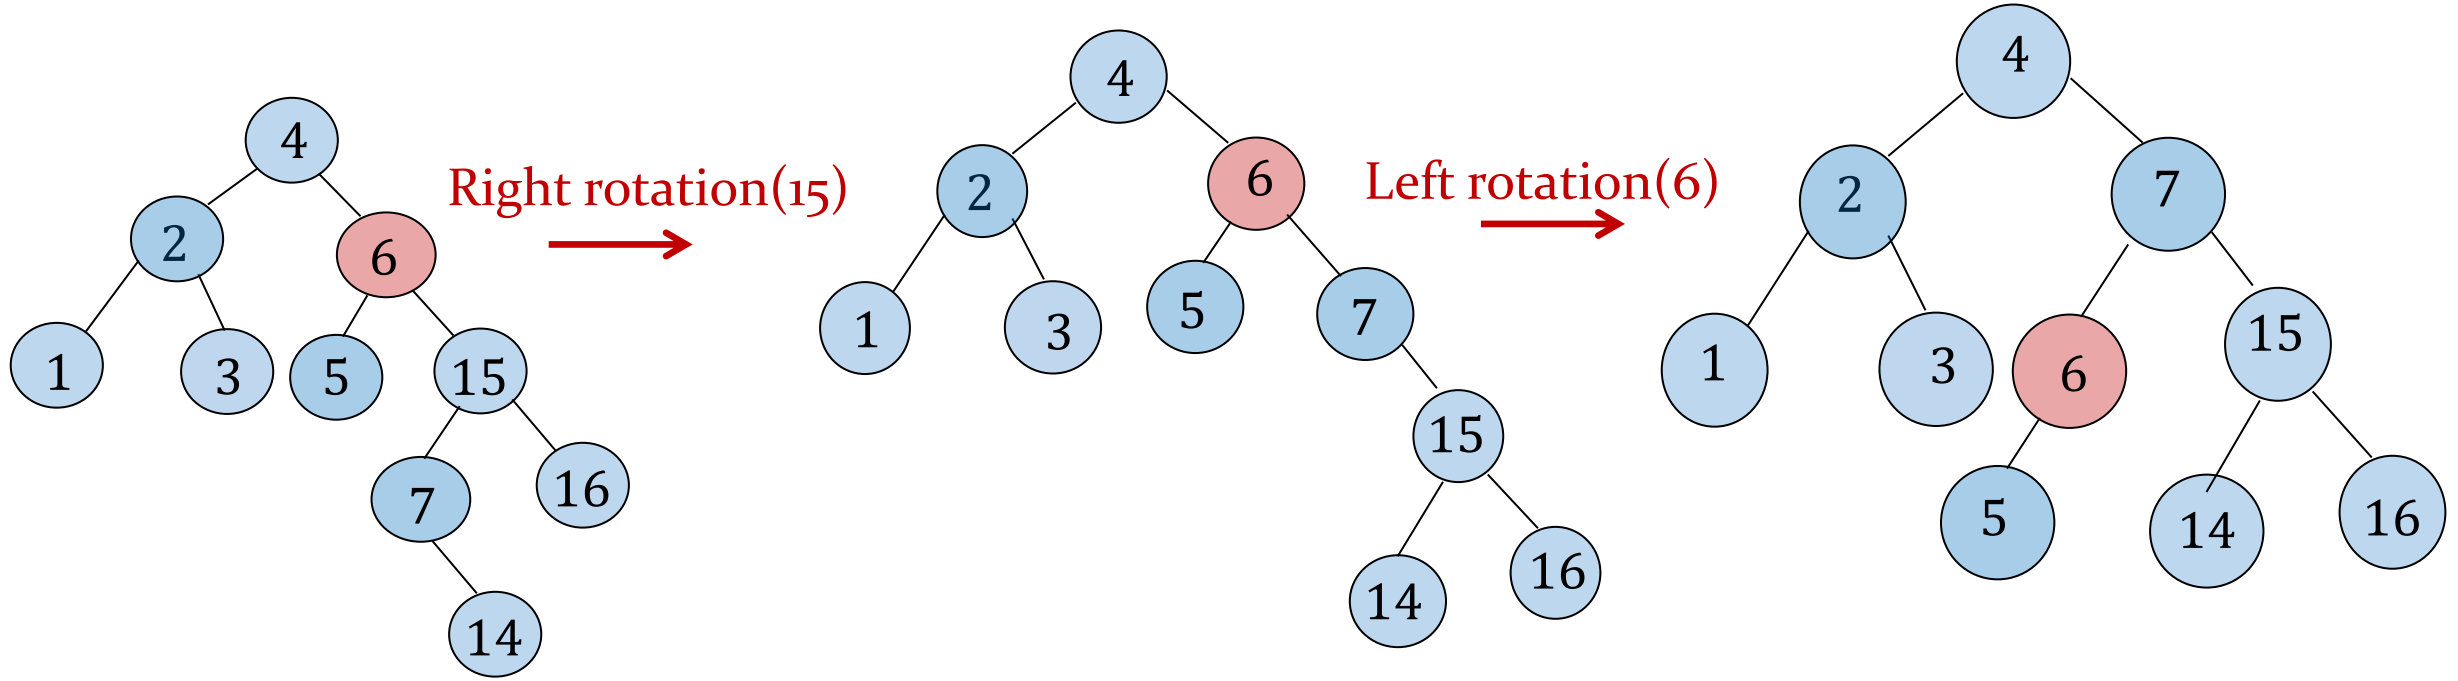


Fig 18

Example - continue

Insert: 3, 2, 1, 4, 5, 6, 7, 16, 15, 14



Example - continue

Insert: 3, 2, 1, 4, 5, 6, 7, 16, 15, 14




Fig 19

Left rotation(6)
→





Fig 20

Update height

Remark: When a node is inserted only the height of its parent and ancestors will be updated.



AVL-Rebalance – pseudocode

AVL – REBALANCE – TO – THE – LEFT(root):
 if *root.right.left.height > root.right.right.height*:
 root.right \leftarrow *AVL – ROTATE – TO – THE – RIGHT*(root.right)
return *AVL – ROTATE – TO – THE – LEFT*(root)

AVL – REBALANCE – TO – THE – RIGHT(root):
 if *root.left.right.height > root.left.left.height*:
 root.left \leftarrow *AVL – ROTATE – TO – THE – LEFT*(root.left)
return *AVL – ROTATE – TO – THE – RIGHT*(root)

AVL-Insert – pseudocode

AVL – INSERT(*root*, *x*):

if *root* is *NIL*:

 # found insertion point

root \leftarrow *TreeNode*(*x*)

 # will be returned below

 # assumption: *TreeNode* creates node with .height = 0

elif *x*.key < *root*.item.key:

root.left \leftarrow *AVL – INSERT*(*root*.left, *x*)

if *root*.left.height > *root*.right.height + 1:

root \leftarrow *AVL – REBALANCE – TO – THE – RIGHT*(*root*)

else:

 # no rebalancing, but height might have changed

AVL – UPDATE – HEIGHT(*root*)

elif *x*.key > *root*.item.key:

root.right \leftarrow *AVL – INSERT*(*root*.right, *x*)

if *root*.right.height > *root*.left.height + 1:

root \leftarrow *AVL – REBALANCE – TO – THE – LEFT*(*root*)

else:

 # no rebalancing, but height might have changed

AVL – UPDATE – HEIGHT(*root*)

else:

 # *x*.key = *root*.item.key Just replace *root*'s item with *x* nothing else changes.

root.item \leftarrow *x*

return *root*

Running time of AVL-Insert

Just Tree-Insert plus some constant time for rotations and height updating.

Overall, worst case $O(h)$ since it's balanced, $O(\log n)$


AVL-Delete($root, x$)

Delete node x from the AVL tree rooted at $root$


Cases that need rebalancing

Height of the “whole subtree” rooted at x before deletion is $h + 3$

Case 1: the deletion reduces the height of a node's **right subtree**, and that node was **left heavy**.



Case 2: the insertion increases the height of a node's **left subtree**, and that node was already **left heavy**.




AVL-Delete: General idea


- First do a normal BST Tree-Delete
- The deletion may cause changes of subtree heights, and may cause certain nodes to lose AVL-ness ($BF(x)$ is 0, 1 or -1)
- Then rebalance by single or double rotations, similar to what we did for AVL-Insert.
- Then update height (BFs) of affected nodes.

Homework: Write the pseudocode for AVL-Delete. (Bring it to the class next week: It is part of your participation)


AVL-Delete (Example)





Delete(13)



AVL-Delete (Example)



Delete(2)



Can we search faster than $O(\log n)$?

Next week we discuss Hash table

Example: Ordered Set

An ADT with the following operation

- **Search**(S, k) in $O(\log n)$
 - **Insert**(S, x) in $O(\log n)$
 - **Delete**(S, x) in $O(\log n)$
-
- **Rank**(k): return the rank of key k
 - **Select**(r): return the key with rank r

E.g., $S = \{ 27, 56, 30, 3, 15 \}$

$\text{Rank}(15) = 2$ because 15 is the second smallest key

$\text{Select}(4) = 30$ because 30 is the 4th smallest key

Augmentation needed

Questions