


Agenda

- FFS (Fast File System)
 - Device aware
- NTFS - Windows extent-based file system
- Supporting multiple file systems - VFS (Virtual File System)
- FYI - exercises marks will show up on **MarkUs** (not Quercus)

The Fast File System:

An example of a device aware
file system

The common storage device interface



OS's view of storage device

- Storage exposed as linear array of blocks
- Common block sizes: 512 bytes, 4096 bytes
- Number of blocks: device capacity / block size

Back to file systems

- Key idea: File systems need to be aware of disk characteristics for performance
 - Allocation algorithms to enhance performance
 - Request scheduling to reduce seek time


Enhancing disk performance

- High-level disk characteristics yield two goals:
 - **Closeness**
 - reduce seek times by putting related things close to one another
 - generally, benefits can be in the factor of 2 range
 - **Amortization**
 - amortize each positioning delay by grabbing lots of useful data
 - generally, benefits can reach into the factor of 10 range

Allocation Strategies

- Disks perform best if seeks are reduced and large transfers are used
- **Scheduling requests** is one way to achieve this
- **Allocating related data “close together”** on the disk is even more important

Inodes: Indirection & Independence




- + File size grows dynamically, allocations are independent
- Hard to achieve closeness and amortization

FFS: A disk-aware file
system

Original Unix File System


- Recall FS sees storage as linear array of blocks
- Each block has a logical block number (LBN)




- Simple, straightforward implementation
 - Easy to implement and understand
 - But very poor utilization of disk bandwidth (lots of seeking)

Data and Inode Placement: Problem 1


- On a new FS, blocks are allocated sequentially, close to each other.



- As the FS gets older, files are being deleted and create random gaps



- In aging file systems, data blocks end up allocated far from each other:



- Data blocks for new files end up scattered across the disk!
- Fragmentation of an aging file system causes more seeking!

Data and Inode Placement – problem #2

Superblock



- Inodes allocated far from blocks
 - All inodes at beginning of disk, far from data
- Recall that when we traverse a file path, at each level we inspect the inode first, then access the data block.
 - Traversing file name paths, manipulating files, directories requires **going back and forth from inodes to data blocks**
- => Again, lots of seeks!


FFS

- BSD Unix folks did a redesign (BSD 4.2) that they called the Fast File System (FFS)
 - Improved disk utilization, decreased response time
 - McKusick, Joy, Leffler, and Fabry, ACM TOCS, Aug. 1984
- Now the FS from which all other Unix FS's have been compared
- Good example of being device-aware for performance



Cylinder Groups

- BSD FFS addressed placement problems using the notion of a cylinder group (aka allocation groups in lots of modern FS's)
 - Disk partitioned into groups of cylinders
 - Data blocks in same file allocated in same cylinder group
 - Files in same directory allocated in same cylinder group
 - Inodes for files allocated in same cylinder group as file data blocks



Cylinder group organization

Cylinder Groups (cont'd)

- Allocation in cylinder groups provides closeness
 - Reduces number of long seeks
- Free space requirement
 - To be able to allocate according to cylinder groups, the disk must have free space scattered across cylinders
 - 10% of the disk is reserved just for this purpose
 - When allocating a large file, break it into large chunks and allocate from different cylinder groups, so it does not fill up one cylinder group
 - If preferred cylinder group is full, allocate from a “nearby” group

More FFS solutions

- Small blocks (1K) in orig. Unix FS caused 2 problems:
 - Low bandwidth utilization
 - Small max file size (function of block size)
- Fix using a larger block (4K)
 - Very large files, only need two levels of indirection for 2^{32}
 - New Problem: internal fragmentation
 - Fix: Introduce “fragments” (1K pieces of a block)
- Problem: Media failures
 - Replicate master block (superblock)
- Problem: Device oblivious
 - Parameterize according to device characteristics

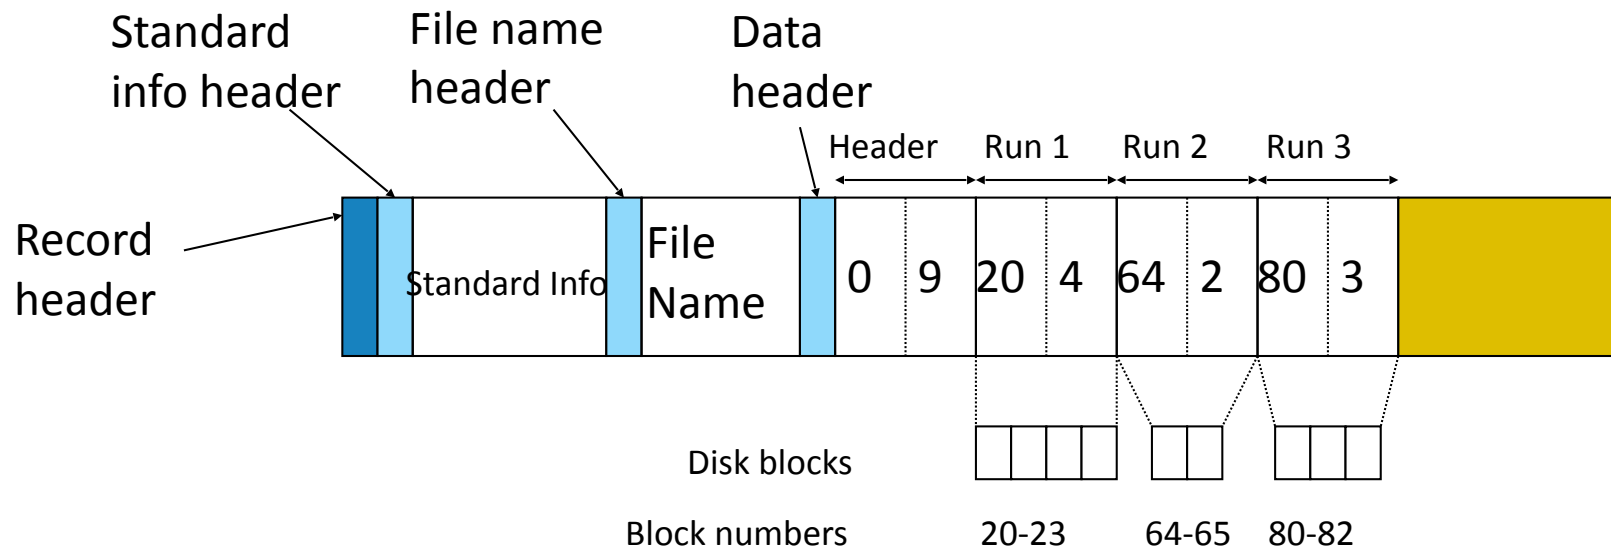
NTFS

- The New Technology File System (NTFS) from Microsoft replaced the old FAT file system.
- The designers had the following goals:
 1. Eliminate fixed-size short names
 2. Implement a more thorough permissions scheme
 3. Provide good performance
 4. Support large files
 5. Provide extra functionality:
 - Compression
 - Encryption
 - Types
- In other words, they wanted a file system flexible enough to support future needs.

NTFS

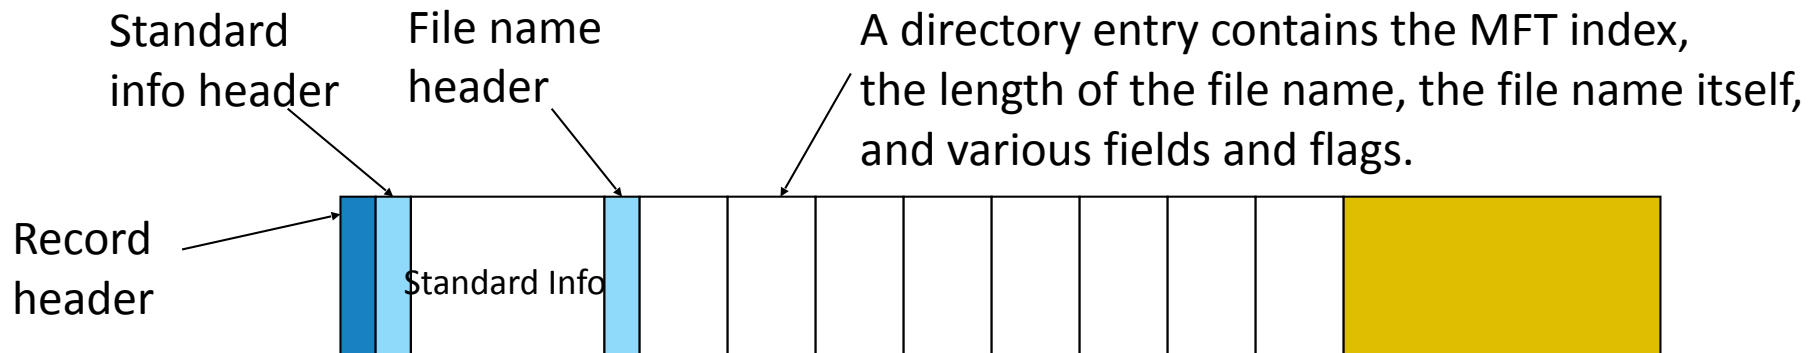
- Each volume (partition) is a linear sequence of blocks (usually 4 Kb block size).
- Each volume has a Master File Table (MFT).
 - Sequence of 1 KB records.
 - One or more record per file or directory
 - Similar to inodes, but more flexible
 - Each MFT record is a sequence of variable length (attribute header, value) pairs.
 - Long attributes can be stored externally, and a pointer kept in the MFT record.
- NTFS tries to allocate files in runs of consecutive blocks.

MFT Record




- An MFT record for a 3-run 9-block file.
- Each “data” attribute indicates the starting block and the number of blocks in a “run” (or extent)
- If all the records don’t fit into one MFT record, extension records can be used to hold more.

MFT Record for a Small Directory



- Directory entries are stored as a simple list
- Large directories use B+ trees instead.

MFT Small file



- For very small files, data can be stored in the MFT record

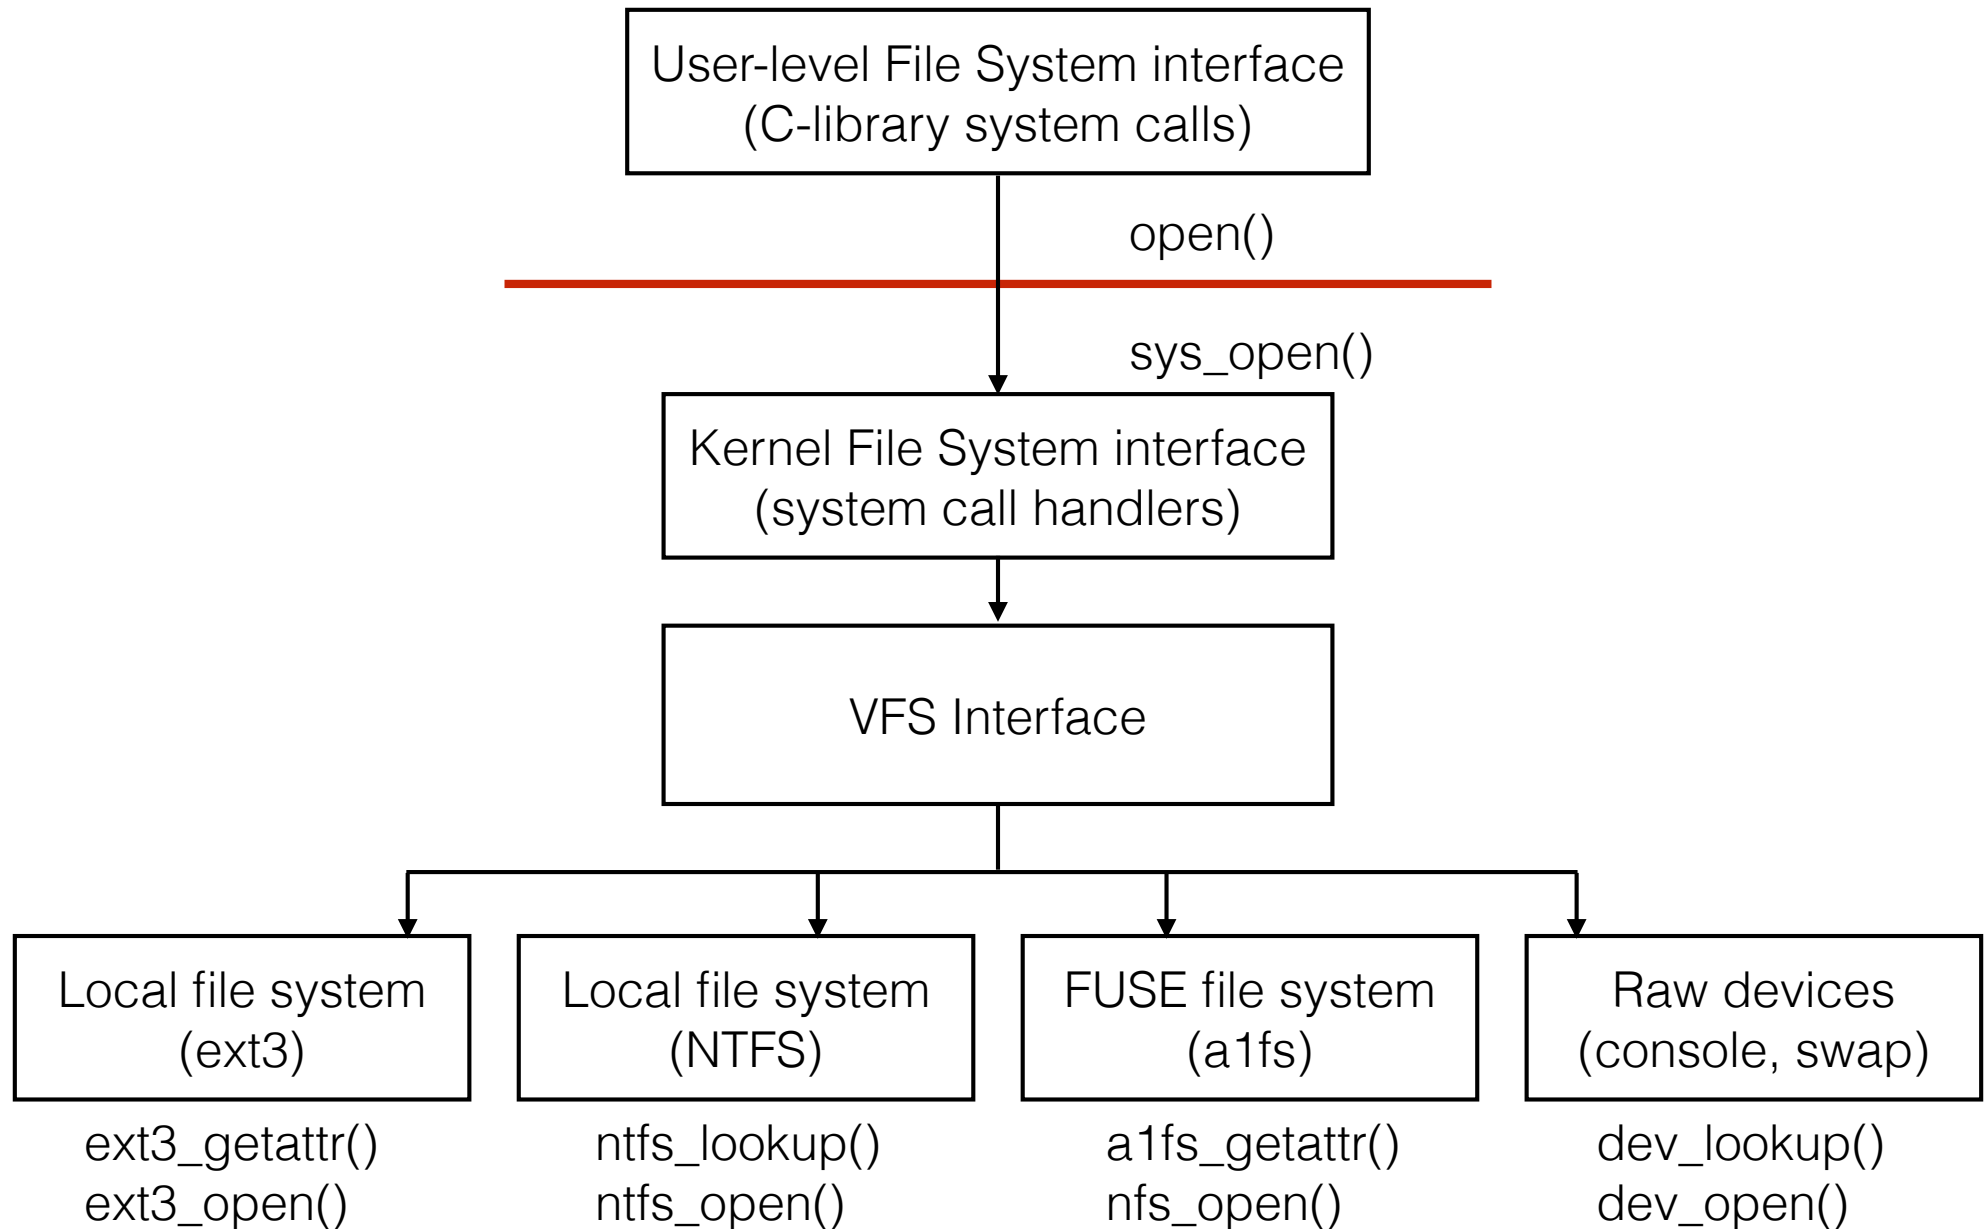
NTFS

- Metadata (attributes)
 - key-value pairs
 - significant flexibility
 - allows implementation of extra features:
compression, different file types

Ext2, Ext3, Ext4

- Linux file system evolution
- Ext2 originally borrowed heavily from FFS
- Recall: Reduce seeks for faster reads, etc
- More details on reliability and optimizations for writes

Schematic View of VFS



Supporting Multiple File Systems

VFS (Virtual File System)

- Provides an abstract file system interface
 - Separates abstraction of file and collections of files from specific implementations
 - System calls such as open, read, write, etc. can be implemented in terms of operations on the abstract file system
 - `vfs_open`, `vfs_close`
- Abstraction layer is for the OS itself
 - user-level programmer interacts with the file systems through the system calls