


Reliability and Write Optimizations

Reliability and Write Optimizations

- How do we guarantee consistency of on-disk storage?
- How do we handle OS crashes and disk errors?
- How do we optimize writes?

FFS: Consistency Issues - Overview



- Inodes: fixed size structure stored in cylinder groups
- Metadata updates must be synchronous operations.
Why?
- File system operations affect multiple metadata blocks
 - Write newly allocated inode to disk before its name is entered in a directory.
 - Remove a directory name before the inode is deallocated
 - Deallocate an inode (mark as free in bitmap) before that file's data blocks are placed into the cylinder group free list.

FFS Observation 1: Crash recovery

- If the OS crashes in between any of these synchronous operations, then the file system is in an inconsistent state.
- Solutions (overview):
 - **fsck** – post-crash recovery process to scan file system structure and restore consistency
 - All data blocks pointed to by inodes (and indirect blocks) must be marked allocated in the data bitmap
 - All allocated inodes must be in some dir entries
 - Inode link count must match
 - Log updates to enable roll-back or roll-forward

Example: update

- Consider a simple update: append 1 data block to a file
- Assume a similar FS structure as seen before:

Inode Bitmap				Data Bitmap				Inodes				Data blocks			
0	1	0	0	0	0	0	0		I[v1]					Da	
0	0	0	0	1	0	0	0								

- Add a data block: Db .. What changes?

Inode Bitmap				Data Bitmap				Inodes				Data blocks			
0	1	0	0	0	0	0	0		I[v2]					Da	Db
0	0	0	0	1	1	0	0								

- Three writes: I[v2], Data Bitmap, Db

Example2: create file

- Create a new **empty** file in the first block of the directory that Inode 1 represents

Inode Bitmap				Data Bitmap				Inodes				Data blocks			
0	1	0	0	0	0	1	1		I[v2]					Da	Db
0	0	0	0	0	0	0	0								

Inode Bitmap				Data Bitmap				Inodes				Data blocks			
0	1	1	0	0	0	1	1		I[v3]	I[v1]				Da	Db
0	0	0	0	0	0	0	0								

- Writes: I[v3], I[v1], Inode Bitmap, Data block

Crash Consistency

- What if only one write succeeds before a crash?

1. Just Db write succeeds

2. Just I[v2] write succeeds

3. Just B[v2] write succeeds

Other Crash Scenarios

- What if only two writes succeed before a crash?
 1. Only $I[v2]$ and Data Bitmap writes succeed.
 2. Only $I[v2]$ and Db writes succeed.
 3. Only Data Bitmap and Db writes succeed.

Solution #1 fsck

- fsck: UNIX tool for finding inconsistencies and repairing them
- Cannot fix all problems!
 - When Db is garbage – cannot know that's the case
 - Only cares that FS metadata is consistent!
- Similar tools exist on other systems

fsck

- What does it check?
 - 1. Superblock: sanity checks
 - Use another superblock copy if suspected corruption
 - 2. Free blocks: scan inodes (incl. all indirect blocks), build bitmap
 - inodes / data bitmaps inconsistency => resolve by trusting inodes
 - Ensure inodes in use are marked in inode bitmaps
 - 3. Inode state: check inode fields for possible corruption
 - e.g., must have a valid “mode” field (file, dir, link, etc.)
 - If cannot fix => remove inode and update inode bitmap
 - 4. Inode links: verify links# for each inode
 - Traverse directory tree, compute expected links#, fix if needed
 - If inode discovered, but no dir refers to it => move to “lost+found”

5. Duplicates: check if two different inodes refer to same block

- Clear one if obviously bad, or, give each inode its own copy of block

6. Bad blocks: bad pointers (outside of valid range)

- Just remove the pointer from the inode or indirect block

7. Directory checks: integrity of directory structure

- E.g., make sure that “.” and “..” are the first entries, each inode in a directory entry is allocated, no directory is linked more than once

fsck limitations

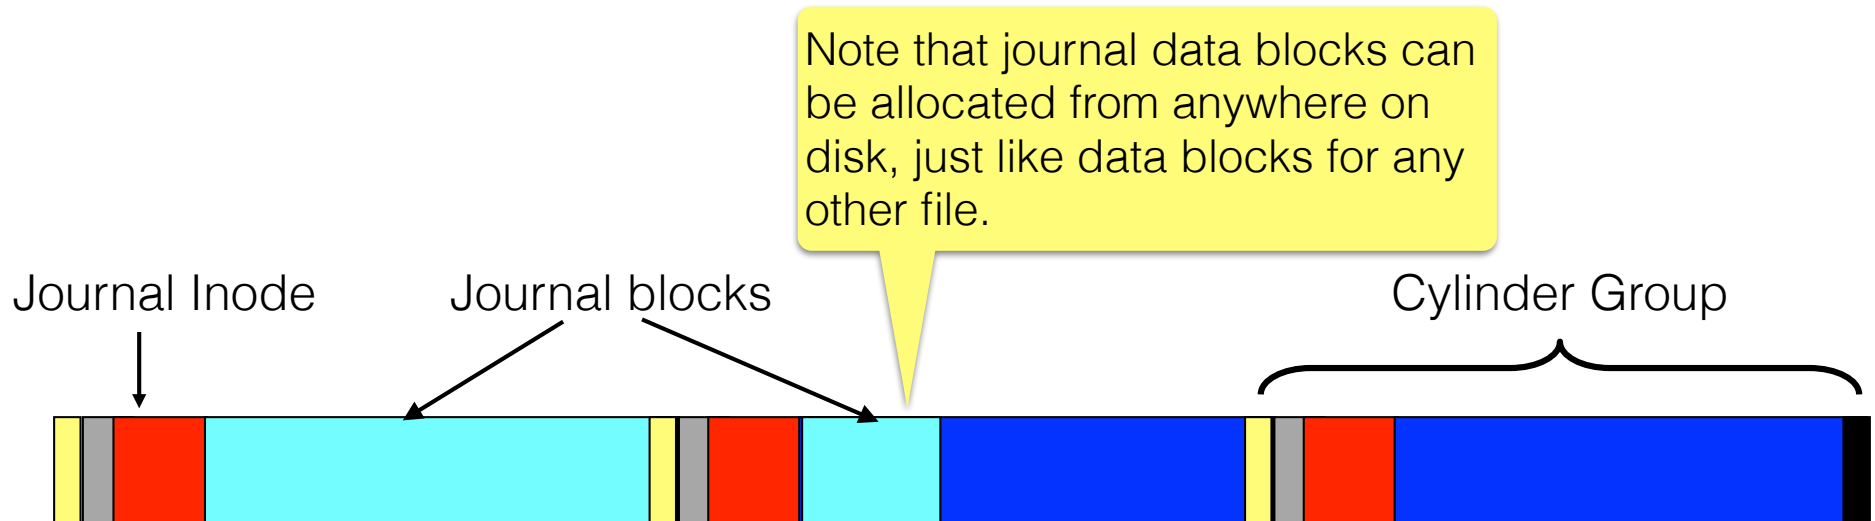
- So, fsck helps ensure integrity
- Only FS integrity, cannot do anything about lost data!
- Bigger problem: **too slow!**
 - Disks are very large nowadays – scanning all this could take hours!
 - Even for small inconsistency, must scan whole disk!

Alternative solution: Journaling

- Aka Write-Ahead-Logging
- Basic idea:
 - When doing an update, before overwriting structures, first write down a little note (elsewhere on disk) saying what you plan to do.
 - i.e., “Log” the operations you are about to do.
- If a crash takes place during the actual write => go back to journal and retry the actual writes.
 - Don't need to scan the entire disk, we know what to do!
 - Can recover data as well
- If a crash happens before journal write finishes, then it doesn't matter since the actual write has NOT happened at all, so nothing is inconsistent.

Linux Ext3 File System

- Extends ext2 with journaling capabilities
 - Backwards and forwards compatible
 - Identical on-disk format
 - Journal can be just another large file (inode, indirect blocks, data blocks)



What goes in the “log”


- Transaction structure:
 - Starts with a “transaction begin” (TxBegin) block, containing a transaction ID
 - Followed by blocks with the content to be written
 - Physical logging: log exact physical content
 - Logical logging: log more compact logical representation
 - Ends with a “transaction end” (TxEnd) block, containing the corresponding TID

Journal
Entry

TxBegin (TID=1)	Updated inode	Updated Bitmap	Updated Data block	TxEnd (TID=1)
--------------------	------------------	-------------------	-----------------------	------------------


Data Journaling Example

- Say we have a regular update – add 1 data block to a file:
 - Write inode (I[v2]), Bitmap (B[v2]), Data block (Db)
 - Markers for the log (transaction begin/end)



Data Journaling Example


- Say we have a regular update – add 1 data block to a file:
 - Write inode ($I[v2]$), Bitmap ($B[v2]$), Data block (Db)
 - Markers for the log (transaction begin/end)




- Sequence of operations
 1. Write the transaction (containing $Iv2$, $Bv2$, Db) to the log
 2. Write the blocks ($Iv2$, $Bv2$, Db) to the file system
 3. Mark the transaction free in the journal
- Crash may happen at any point!
 - If between 1 and 2 => on reboot, replay non-free transactions (called redo logging)
 - If during writes to the journal (step 1) => tricky!

Data Journaling Example

- One solution: write each block at a time
 - Slow!
 - Ideally issue multiple blocks at once.
 - Unsafe though! What could happen?
 - Normal operation: Blocks get written in order, power cuts off before TxEnd gets written => We know transaction is not valid, no problem.



- However, Internal disk scheduling: TxBegin, Iv2, Bv2, TxEnd, Db
- Disk may lose power before Db written




- Problem: Looks like a valid transaction!


Data Journaling Example

To avoid this, split into 2 steps

1. Write all except TxEnd to journal (**Journal Write step**)



2. Write TxEnd (only once 1. completes) (**Journal Commit step**)
=> final state is safe!



3. Finally, now that journal entry is safe, write the actual data and metadata to their right locations on the FS (**Checkpoint step**)
4. Mark transaction as free in journal (**Free step**)

Journaling: Recovery Summary

- If crash happens before the transaction is committed to the journal
 - Just skip the pending update
- If crash happens during the checkpoint step
 - After reboot, scan the journal and look for committed transactions
 - Replay these transactions
 - After replay, the FS is guaranteed to be consistent
 - Called redo logging

Journal Space Requirements

- How much space do we need for the journal?
 - For every update, we log to the journal => sounds like it's huge!
- After “checkpoint” step, the transaction is not needed anymore because metadata and data made it safely to disk
 - So the space can be freed (free step).
- In practice: **circular log**

Metadata Journaling

- Recovery is much faster with journaling
 - Replay only a few transactions instead of checking the whole disk
- However, normal operations are slower
 - Every update must write to the journal first, then do the update
 - Writing time is at least doubled
 - Journal writing may break sequential writing. Why?
 - Jump back-and-forth between writes to journal and writes to main region
 - Metadata journaling is similar, except we only write FS metadata (no actual data) to the journal:

Journal
Entry

TxBegin	I[v2]	B[v2]	TxEnd
---------	-------	-------	-------

Metadata Journaling

- What can happen now?
 - Say we write data after checkpointing metadata
 - If crash occurs before all data is written, inodes will point to garbage data!
 - How do we take care of this?
- Write data BEFORE writing metadata to journal!
 1. Write data, wait until it completes
 2. Metadata journal write
 3. Metadata journal commit
 - 4.4. Checkpoint metadata
 5. Free
- If write data fails => as if nothing happened, sort of (from the FS's point of view)!
- If write metadata fails => same!

Summary: Journaling

- Journaling ensures file system consistency
- Complexity is in the size of the journal, not the size of the disk!
- Is fsck useless then?
- Metadata journaling is the most commonly used
 - Reduces the amount of traffic to the journal, and provides reasonable consistency guarantees at the same time.
- Widely adopted in most modern file systems (ext3, ext4, ReiserFS, JFS, XFS, NTFS, etc.)

Ext3 final notes

- Lacks modern FS features (e.g., extents)
 - For recoverability, this may actually be an advantage
 - FS metadata is in fixed, well-known locations, and data structures have redundancy
 - When faced with significant data corruption, ext2/3 may be recoverable when a tree-based file-system may not

Next up...

- Log-structured file systems

FFS

- Disk block index stored in inodes
- Metadata stored in inodes
- Directory entry stores file name and inode number
- Free blocks: bitmap
- Read performance?
 - Locate related blocks in same cylinder group
 - Locate inodes close to data blocks
- Write performance?
 - Block reallocation – reduces fragmentation, controls aging
 - Soft updates – alternative to journaling; ensures consistency without limiting performance
 - For rest of failure issues – background fsck


FFS Observation 2: Performance

- Performance of FFS: optimized for disk block clustering, using properties of the disk to inform file system layout
- Observation: Memory is now large enough
 - Most reads that go to the disk are the first read of a file. Subsequent reads are satisfied in memory by file buffer cache.
- => there is no performance problem with reads. But write calls could be made faster.
- Writes are not well-clustered. Why?

Log-Structured File System (LFS)

- Traditional FSs:

- Files laid out with spatial locality in mind



- Changes in place to mitigate seeks => e.g. A1', B2', C2'




- Avoids fragmenting files, keeps locality => Reads perform well

Log-Structured File System (LFS)

- Another approach:
 - Memory is increasing => don't care about reads, most will hit in mem.
 - Assume writes will pose the bigger I/O penalty
=> Treat storage as a circular log (Log Structured File System)
- Positive side-effects?
 - Write throughput improved (batched into large sequential chunks)
 - Crash recovery - simpler
- Disadvantages?
 - Initial assumption may not hold => reads much slower on HDDs. Why?

Log-Structured File System (LFS)

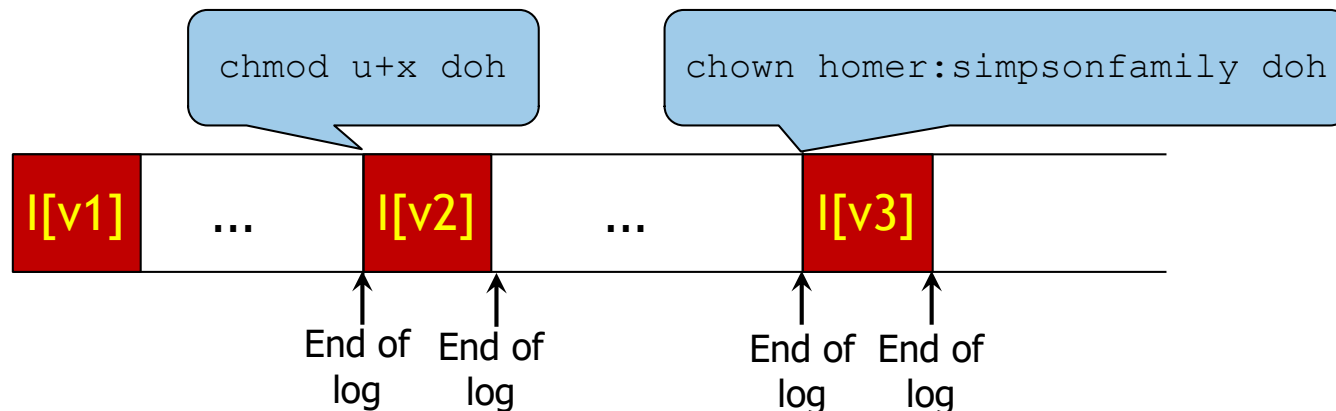
- Ousterhout 1989
- Write all file system data in a continuous log.
- Uses inodes and directories from FFS




- Segments: each has a summary block
- Summary block contains pointer to next one
- Need a fresh segment => first clean an existing partially-used segment (garbage collection)

LFS: Locating inodes

- So, how do we find the inodes though?
 - In typical UNIX FSs, it's easy – array on disk at fixed locations
 - Superblock => Inode Table addr; Then add $\text{Inode\#} * \text{InodeSize}$
- LFS: not so easy. Why?
 - Updates are sequential => inodes scattered all over the disk
 - Also, inodes not static, keep moving




LFS: Locating inodes





- LFS: Needs an inode map (imap) to find the inodes
 - An inode number is no longer a simple index from the start of an inode table as before
- Inode map must keep persistent, to know where inodes are
 - So it has to be on disk as well..
 - So, .. where exactly is the inode map stored on disk?

LFS: Locating inodes

- Put it on a fixed part of the disk
 - Inode map gets updated frequently though
=> Seeks↑ Performance↓




- Instead, place chunks of inode map next to new information



- Hold on, but then how do we now find the imap?

LFS: Locating inodes

- Ok fine! The file system must have some fixed and known location on disk to begin a file lookup
- Checkpoint region (CR)
 - Pointers to the latest pieces of the inode map
 - So, find imap pieces by reading the CR!



Crash Recovery

- What if the system crashes while LFS is writing to disk?
- LFS normally buffers writes in segment
 - When full (or at periodic time intervals), writes segment to disk
- LFS writes segments and periodically also updates the CR
- Crashes can happen at any point
 - How do we handle crashes during these writes?
- Solution:
 1. Uncommitted segments – easy: reconstruct from log after reboot
 2. CR: Keep two CRs, at either end of the disk; alternate writes
 - Update protocol (header, body, last block)

Garbage Collection - Complicated!

- LFS repeatedly writes latest version of a file to new locations on disk
- Older versions of files (garbage) are scattered throughout the disk
- Must periodically find these obsolete versions of file data and clean up => free blocks for subsequent writes
- Cleaning done on a segment-by-segment basis
 - Since the segments are large chunks, it avoids the situation of having small “holes” of free space
- Garbage collection in LFS – interesting research problem!
 - Series of papers were published looking at its performance
 - Depending on when cleaning happens there may be significant performance loss.

LFS Summary

- A new approach that prioritizes update performance
- **Gist:** Instead of overwriting in place, append to log and reclaim (garbage collect) obsolete data later.
- **Advantage:**
 - Very efficient writes
- **Disadvantages:**
 - Less efficient reads (more indirection does not solve everything)
 - But assumes most reads hit in memory anyway.
 - Garbage collection is a tricky problem
- LFS inspired the logging features of journaling file systems in use today. E.g., Ext3/Ext4

Redundancy



- We saw how the journal helps recover from a crash during a write, by replaying the log entries of the write operations
- This assumes that the disk is still usable after rebooting
- What if we have disk failures?
- What can we do to prevent data loss?
 - Have more than one copies of the data: Redundancy!



RAID

- Redundant Array of ~~Inexpensive~~ Independent Disks (RAID)
- Reliability strategies:
 - Data duplicated – mirror images, redundant full copy => one disk fails, we have the mirror
 - Data spread out across multiple disks with redundancy => Can recover from a disk failure, by reconstructing the data
- Concepts:
 - **Redundancy/Mirroring**: keep multiple copies of the same block on different drives, just in case a drive fails
 - **Parity information**: XOR each bit from 2 drives, store checksum on 3rd drive
- Multiple RAID levels – we'll only look at a few

Some Standard RAID levels




Striping

- Files are divided across disks
- Improves throughput
- If one drive fails, the whole volume is lost

Mirroring

- Capacity is half
- Any drive can serve a read
- Improved read throughput
- Write throughput is slower
- If one drive fails, no data lost

RAID 5




- Block level striping
- Distributed parity
- A failed disk can be reconstructed from the rest

Solid State Disks (SSD)

- Replace rotating mechanical disks with non-volatile memory
 - Battery-backed RAM
 - NAND flash
- Advantages: faster
- Disadvantages:
 - Expensive
 - Wear-out (flash-based)
- NAND flash storage technology
 - Read / write / **erase!** operations

SSD Characteristics

- Data cannot be modified “in place”
 - No overwrite without erase
- Terminology:
 - Page (unit of read/write), block (unit of erase operation)



Data written in 4KB pages

Data erased in blocks of
typically ≥ 128 pages

- Uniform random access performance!
 - Disks typically have multiple channels so data can be split (striped) across blocks, speeding access time

Writing

- Consider updating a file system block (e.g. a bitmap allocation block in ext2 file system)
 - Find the block containing the target page
 - Read all active pages in the block into controller memory
 - Update target page with new data in controller memory
 - Erase the block (high voltage to set all bits to 1)
 - Write entire block to drive
- Some FS blocks are frequently updated
 - And SSD blocks wear out (limited erase cycles)

SSD Algorithms

- **Wear levelling**
 - Always write to new location
 - Keep a map from logical FS block number to current SSD block and page location
 - Old versions of logically overwritten pages are “stale”
- **Garbage collection**
 - Reclaiming stale pages and creating empty erased blocks
- **RAID 5** (with parity checking) striping across I/O channels to multiple NAND chips

File Systems and SSDs

- Typically, same FSs as for hard disk drives
 - ext4, Btrfs, XFS, JFS and F2FS support SSDs
- No need for the FS to take care of wear-leveling
 - Done internally by the SSD
 - But the **TRIM** operation is used to tell the SSD which blocks are no longer in use. (Otherwise a delete operation doesn't go to disk)
- Some flash file systems (F2FS, JFFS2) help reduce write amplification (esp. for small updates – e.g., FS metadata)
- Other typical HDD features – do we want these?
 - Defragmentation
 - Disk scheduling algorithms

SSD Reliability

- FAST2016 paper “Flash reliability in production”
Google study on:
 - Millions of drive days over 6 years
 - 10 different drive models
 - 3 different flash types: MLC, eMLC and SLC
 - Enterprise and consumer drives
- Full paper: <https://www.usenix.org/conference/fast16/technical-sessions/presentation/schroeder>
- Key points: <http://www.zdnet.com/article/ssd-reliability-in-the-real-world-googles-experience/>

Summary: File System Goals

- Efficiently translate file name into file number using a directory
- Sequential file access performance
- Efficient random access to any file block
- Efficient support for small files (overhead in terms of space and access time)
- Support large files
- Efficient metadata storage and lookup
- Crash recovery

Summary: File System Components

- Index structure to locate each block of a file
- Free space management
- Locality heuristics
- Crash recovery