

**Question 1.** [6 MARKS]

Each of the following statements is false. In one sentence, explain why. (If you disagree that the statement is false, provide your reasoning.)

**Part (a)** [1 MARK] A trap instruction is a privileged operation.

*A user-level process needs to be able to call trap to switch into kernel-mode.*

**Part (b)** [1 MARK]

Locks and semaphores are identical structures.

*Semaphores have a variable associated with them. More than one thread can pass through a sem\_wait call if the value of the semaphore is greater than 1. (Need to say something that demonstrates knowledge of what a semaphore is.)*

**Part (c)** [1 MARK]

If two threads access a shared variable at the same time, there will be a concurrency error.

*It is fine for two threads to load a shared variable at the same time. They will both see the right value. It may also be fine if only one thread is guaranteed to be updating the shared variable, and the timing of when the second thread looks at the value doesn't matter.*

**Part (d)** [1 MARK]

Limited direct execution means that once the OS has prepared a process for execution, the process has full control of the CPU.

*The process still cannot execute privileged instructions. It isn't enough to say that the process can be interrupted, because if the process could run privileged instructions, it could disable interrupts or overwrite the interrupt handler.*

**Part (e)** [1 MARK]

Indexed-based file systems suffer from external fragmentation because blocks of the file may be scattered across the disk.

*The fact that blocks are scattered across the disk does not prevent regions of block from being used. If a file in an indexed-based file system needs a new block, and there is one free, the file can use it. Therefore external fragmentation cannot occur.*

**Part (f)** [1 MARK]

Extent-based file systems may require more disk block accesses than indexed-based file systems for random access because it may need to traverse an extent to get to a particular byte in the file.

*Can do math to find the block so we don't have to read all of the blocks in an extent. (In fact indexing may be worse due to indirect blocks.)*

**Question 2.** [7 MARKS]

**Part (a)** [3 MARKS]

Name 3 fields stored in the data structure used to keep track of processes in the kernel.

*Examples:* `pid, exit code, address space, program counter, stack pointer, registers describing the current state.`

**Part (b)** [2 MARKS]

Describe what state needs to be saved for a process during a context switch.

*Registers:* `stack pointer, program counter, and user registers. Kernel state`

**Part (c)** [2 MARKS]

Give one concrete example of how the operating system virtualizes the hardware resources that we have seen so far in the course. Identify which resource is virtualized and how the operating system achieves this virtualization.

emphThere are several possibilities here. The OS virtualizes the CPU. A process runs without knowing about context switches. `The file API virtualizes the disk.` Some explanation is necessary for full marks.

**Question 3.** [2 MARKS]

Consider the following excerpt from the solution to the reader/writer problem using condition variables:

```
// assume all variable are correctly initialized
// assume writing will only ever hold the values 0 or 1

pthread_mutex_lock(&mutex);
while(writing) {
    pthread_cond_wait(&turn, &mutex);
}
readcount++;
pthread_mutex_unlock(&mutex);
```

Which of the following statements are true?

- ☐ `writing` must be 0 when `pthread_cond_wait` returns.
- ☐ The mutex is held by this thread while it is blocked in `pthread_cond_wait`
- ☒ The mutex is held by this thread whenever it checks the value of `writing`
- ☒ The mutex is held by this thread when the while loop terminates.

**Question 4.** [6 MARKS]

A program issues the following system calls on the original Fast File System (FFS) file system. The file system block size is 4KB and the size of an inode is 64 bytes. Assume that the file system is already mounted when the program starts and that all system calls succeed.

```
char buf[4096];
int fd = open("/a/b/c", 0); // open in read-only mode
lseek(fd, 1034*4096, 0);    // seek to position (1034*4096) from start of file
read(fd, buf, 4096);        // read 4k of data from file
```

Assume that the file buffer cache is empty. State the *minimum* number of the following types of blocks that will be read when the program above is run. Explain your answer for each block type.

1. Inode block(s) *1* - we need 4 inodes (for root dir, dir a, dir b, and file c) but at 64 bytes per inode, they can all fit in 1 block if we are lucky.
  
2. Directory block(s) *3* - each dir has its own data blocks, so we need 1 for each dir (root, a and b)
  
3. Indirect block(s) (include single, double or triple indirect) *1* - Original Unix FS had 12 direct pointers, 4 byte block addresses, so indirect block holds 1024 block ptrs, so blocks 12-1035 can be read with 1 indirect block. We are reading block 1034.
  
4. Other data block(s) *1* - we seek to a block-aligned position and read 1 block of data.

**Question 5.** File system consistency [9 MARKS]

On an ext2 or FFS file system, consider the operation of creating a new directory in an existing directory. Assume the existing directory occupies one block and there is enough space to add a new entry. Assume the existing directory inode and the new directory inode are in different disk blocks.

**Part (a)** [2 MARKS]

Which of the following blocks must be updated? Check all that apply.

- ☒ inode bitmap
- ☒ new directory inode
- ☒ new directory inode
- ☒ new directory data block
- ☒ existing directory inode
- ☒ existing directory data block

**Part (b)** [2 MARKS]

What data is updated in the existing directory inode?

*last modified time, number of links*

**Part (c)** [5 MARKS]

In each of the remaining questions, check all of the boxes that most closely explain what happens if a crash occurs after updating only the block(s) specified.

**Inode Bitmap and Data Block Bitmap**

- ☐ No inconsistency (it simply appears that the operation was not performed)
- ☒ Data leak (data block is lost for any future use)
- ☒ Inode leak (Inode is lost for any future use)
- ☐ Inconsistent inode data (Some inode field does not match what is stored in data blocks)
- ☐ Multiple file paths may point to same inode
- ☐ Something points to garbage

**New Directory Inode**

- ☒ No inconsistency (it simply appears that the operation was not performed)
- ☐ Data leak (data block is lost for any future use)
- ☐ Inode leak (Inode is lost for any future use)
- ☐ Multiple file paths may point to same inode
- ☐ Inconsistent inode data (Some inode field does not match what is stored in data blocks)

- ☐ Something points to garbage

**Inode Bitmap, Data Block Bitmap, Existing Directory data, New Directory inode, and New Directory data**

**Existing inode: last modified time, number of links**

- ☐ No inconsistency (it simply appears that the operation was not performed)
- ☐ Data leak (data block is lost for any future use)
- ☐ Inode leak (Inode is lost for any future use)
- ☐ Multiple file paths may point to same inode
- ☒ Inconsistent inode data (Some inode field does not match what is stored in data blocks)
- ☐ Something points to garbage

**Inode Bitmap and New Directory inode**

- ☐ No inconsistency (it simply appears that the operation was not performed)
- ☐ Data leak (data block is lost for any future use)
- ☒ Inode leak (Inode is lost for any future use)
- ☐ Multiple file paths may point to same inode
- ☐ Inconsistent inode data (Some inode field does not match what is stored in data blocks)
- ☐ Something points to garbage

**New Directory inode, Existing Directory inode, and Existing Directory data**

- ☐ No inconsistency (it simply appears that the operation was not performed)
- ☐ Data leak (data block is lost for any future use)
- ☐ Inode leak (Inode is lost for any future use)
- ☒ Multiple file paths may point to same inode
- ☐ Inconsistent inode data (Some inode field does not match what is stored in data blocks)
- ☒ Something points to garbage