

1. (a) Order the requests by non-increasing client time, *i.e.*, so that $c_{i_1} \geq c_{i_2} \geq \dots \geq c_{i_n}$.

This takes time $\Theta(n \log n)$, for sorting.

- (b) We introduce some notation to make the proof easier to write. For any permutation $S = i_1, i_2, \dots, i_n$ of $[1, 2, \dots, n]$ and for $1 \leq k \leq n$, let $T_k(S) = s_{i_1} + s_{i_2} + \dots + s_{i_k} + c_{i_k}$ denote the completion time for the k -th request. The overall completion time is then $T(S) = \max \{T_1(S), T_2(S), \dots, T_n(S)\}$.

To prove that the algorithm produces an optimal ordering, we show that every initial portion of the solution (i_1, i_2, \dots, i_k) can be extended to an optimal solution, by induction on k .

Base Case: The empty sequence is extended by all optimal solutions.

Ind. Hyp.: Suppose $k \geq 0$ and i_1, \dots, i_k can be extended to some optimal $S^* = i_1, \dots, i_k, i_{k+1}^*, \dots, i_n^*$.

Ind. Step: Consider the greedy choice for the next index, i_{k+1} . If $i_{k+1} = i_{k+1}^*$ then we are done: S^* already extends i_1, \dots, i_k, i_{k+1} .

Else, let j be such that $i_j^* = i_{k+1}$. Since S^* starts with i_1, \dots, i_k , and $i_{k+1}^* \neq i_{k+1}$, we know that $k+2 \leq j \leq n$. Because of the way i_{k+1} is chosen, we also know that $c_{i_{k+1}} \geq c_{i_{k+1}^*}, c_{i_{k+1}} \geq c_{i_{k+2}^*}, \dots, c_{i_{k+1}} \geq c_{i_n^*}$.

Let $S' = i_1, \dots, i_k, i_{k+1}, i_{k+2}^*, \dots, i_{j-1}^*, i_{k+1}^*, i_{j+1}^*, \dots, i_n^*$ (*i.e.*, S' is the same as S^* except that $i_{k+1} = i_j^*$ and i_{k+1}^* have been swapped). To compare $T(S')$ and $T(S^*)$, we examine the completion time of individual requests:

- $T_1(S') = T_1(S^*), \dots, T_k(S') = T_k(S^*)$ because both S' and S^* start with the same i_1, \dots, i_k .
- $T_{k+1}(S') = s_{i_1} + \dots + s_{i_k} + s_{i_{k+1}} + c_{i_{k+1}}$
 $\leq s_{i_1} + \dots + s_{i_k} + s_{i_{k+1}^*} + s_{i_{k+2}^*} + \dots + s_{i_{j-1}^*} + s_{i_{k+1}} + c_{i_{k+1}}$
 $= T_j(S^*)$
 $T_{k+1}(S^*) = s_{i_1} + \dots + s_{i_k} + s_{i_{k+1}^*} + c_{i_{k+1}^*}$
 $\leq s_{i_1} + \dots + s_{i_k} + s_{i_{k+1}^*} + s_{i_{k+2}^*} + \dots + s_{i_{j-1}^*} + s_{i_{k+1}} + c_{i_{k+1}^*}$
 $\leq s_{i_1} + \dots + s_{i_k} + s_{i_{k+1}^*} + s_{i_{k+2}^*} + \dots + s_{i_{j-1}^*} + s_{i_{k+1}} + c_{i_{k+1}}$
 $= T_j(S^*)$
- Similarly, $T_{k+2}(S') \leq T_j(S^*), \dots, T_{j-1}(S') \leq T_j(S^*)$ and $T_{k+2}(S^*) \leq T_j(S^*), \dots, T_{j-1}(S^*) \leq T_j(S^*)$.
- $T_j(S') = s_{i_1} + \dots + s_{i_k} + s_{i_{k+1}} + s_{i_{k+2}^*} + \dots + s_{i_{j-1}^*} + s_{i_{k+1}^*} + c_{i_{k+1}^*}$
 $= s_{i_1} + \dots + s_{i_k} + s_{i_{k+1}^*} + s_{i_{k+2}^*} + \dots + s_{i_{j-1}^*} + s_{i_{k+1}} + c_{i_{k+1}^*}$
 $\leq s_{i_1} + \dots + s_{i_k} + s_{i_{k+1}^*} + s_{i_{k+2}^*} + \dots + s_{i_{j-1}^*} + s_{i_{k+1}} + c_{i_{k+1}}$
 $= T_j(S^*)$
- $T_{j+1}(S') = T_{j+1}(S^*), \dots, T_n(S') = T_n(S^*)$ because both S' and S^* contain the same requests after request number j .

Hence,

$$\begin{aligned} T(S') &= \max \{T_1(S'), \dots, T_k(S'), T_{k+1}(S'), T_{k+2}(S'), \dots, T_{j-1}(S'), T_j(S'), T_{j+1}(S'), \dots, T_n(S')\} \\ &\leq \max \{T_1(S'), \dots, T_k(S'), T_j(S^*), T_{j+1}(S'), \dots, T_n(S')\} \\ &= \max \{T_1(S^*), \dots, T_k(S^*), T_j(S^*), T_{j+1}(S^*), \dots, T_n(S^*)\} \\ &= \max \{T_1(S^*), \dots, T_k(S^*), T_{k+1}(S^*), T_{k+2}(S^*), \dots, T_{j-1}(S^*), T_j(S^*), T_{j+1}(S^*), \dots, T_n(S^*)\} \\ &= T(S^*). \end{aligned}$$

Since S^* was optimal, this means $T(S') = T(S^*)$. So S' is optimal and extends i_1, \dots, i_{k+1} .

Since every initial portion of the greedy permutation can be extended to an optimal solution, in particular, the final permutation itself is optimal.

2. (a) Counter-example: $G = (V, E)$ with $V = \{a, b, c\}$, $E = \{(a, b), (b, c)\}$, $c(a, b) = c(b, c) = 1$, and $L = \{b\}$. Since G is already a tree, there is only one spanning tree of G (G itself), and b is not a leaf in this tree.
- (b) Consider the graph $G = (V, E)$ with $V = \{a, b, c\}$, $E = \{(a, b), (b, c), (c, a)\}$, $c(a, b) = c(b, c) = 1$, $c(a, c) = 2$, and $L = \{b\}$. Then G contains exactly one MST: $T = \{(a, b), (b, c)\}$, but T is not a solution to the MST with Fixed Leaves problem because b is not a leaf in T .

With the same input, there are two optimal solutions to the MST with Fixed Leaves problem: $T_1 = \{(a, c), (a, b)\}$ and $T_2 = \{(a, c), (b, c)\}$. Neither of these is a MST in G .

- (c) # This is a variation of Kruskal's algorithm: we construct the tree edge-by-edge,
 # starting with the nodes in L .
 # Handle the only special case when two nodes of L must be connected to each other.
if $V = \{a, b\}, E = \{(a, b)\}, L = \{a, b\}$: **return** $\{(a, b)\}$
 # Now, handle the general case.
 $T \leftarrow \emptyset$
 # Start by selecting one edge (v, u) for each node $v \in L$, where $u \notin L$ and $c(v, u)$ is minimum.
for $v \in L$:
 $c \leftarrow \infty$
 for $(v, u) \in E$:
 # Remove all edges adjacent to v from E , to ensure v is a leaf in T .
 $E \leftarrow E - \{(v, u)\}$
 if $u \notin L$ **and** $c(v, u) < c$:
 $e \leftarrow (v, u)$
 $c \leftarrow c(v, u)$
 # At this point, e is a minimum-cost edge connecting v to $V - L$.
 $T \leftarrow T \cup \{e\}$
 Now, run Kruskal's algorithm on the remaining graph, starting from the edges already in T .
return T

The algorithm's complexity is $\mathcal{O}(m \log m)$, same as Kruskal's, since that is the part of the algorithm that takes the longest.

- (d) Let $G_0 = G - L$ (G_0 is G with every node of L removed, and every edge that contains a node from L removed).

Consider any optimal solution T to the MST with Fixed Leaves problem on input G, L . Now consider $T' = T - \{(v, u) : v \in L\}$ (T' is T with its leaves removed).

Claim: T' is a MST in G_0 .

Proof: For a contradiction, suppose T^* is a spanning tree of G_0 and $c(T^*) < c(T')$. Then $T^* \cup \{(v, u) \in T : v \in L\}$ forms a spanning tree of the original graph G whose total cost is smaller than that of T and where each node of L is a leaf. This contradicts the fact that T is an optimal solution for the original input.

Claim: The spanning tree generated by our algorithm is an optimal solution to the MST with Fixed Leaves problem on input G, L .

Proof: In every optimal solution, each node of L must be connected to $V - L$ by exactly one edge. Any choice other than a minimum-cost edge would increase the cost of the resulting spanning tree and would be sub-optimal. Also, connecting one node of L to another node of L would make it impossible for both nodes to be connected to the rest of the graph and still be leaves (unless G consists of exactly one edge).

Once these edges have been selected, and other edges to the vertices of L eliminated (to guarantee every node of L is a leaf), the remaining graph is simply G_0 and Kruskal's algorithm is guaranteed to find a MST of G_0 . Thus, there is no spanning tree of G where each node of L is a leaf and with a smaller cost.

3. (a) **Greedy strategy:** Repeatedly pick the largest coin remaining.

Counter-example: Consider input $A = 30, c_1 = 25, c_2 = 10, c_3 = 10, c_4 = 10$. The greedy algorithm would pick $c_1 = 25$ and be unable to finish making change, even though $\{c_2, c_3, c_4\}$ is a solution.

- (b) **Step 0:** Describe the recursive structure of sub-problems.

For every optimal solution $i_1 < i_2 < \dots < i_k$, either $i_k = m$ or $i_k < m$.

If $i_k = m$, then $\{i_1, i_2, \dots, i_{k-1}\}$ must be an optimal solution for input $A - c_m, c_1, \dots, c_{m-1}$ (if there were a better solution for input $A - c_m, c_1, \dots, c_{m-1}$, we could simply add c_m to it and get a better overall solution).

If $i_k < m$, then $\{i_1, i_2, \dots, i_k\}$ must be an optimal solution for input A, c_1, \dots, c_{m-1} (trivially).

Step 1: Define an array that stores optimal values for arbitrary sub-problems.

Define $N[k, a]$ to be the minimum number of coins required to make change for amount a using coins c_1, \dots, c_k , for $0 \leq a \leq A, 0 \leq k \leq m$. (Let $N[k, a] = \infty$ when the problem has no solution for input a, c_1, \dots, c_k .)

Step 2: Give a recurrence relation for the array values. (We include justifications for each case of the recurrence.)

$N[k, 0] = 0$ for $0 \leq k \leq m$ (no coin is necessary to make change for amount 0).

$N[0, a] = \infty$ for $1 \leq a \leq A$ (it is impossible to make change for a positive amount without using any coins).

$N[k, a] = N[k-1, a]$ if $c_k > a$, for $1 \leq k \leq m, 1 \leq a \leq A$ (coin c_k cannot be used to make change for amount a if $c_k > a$).

$N[k, a] = \min\{N[k-1, a], 1 + N[k-1, a - c_k]\}$ if $c_k \leq a$, for $1 \leq k \leq m, 1 \leq a \leq A$ (any optimal solution either does not use c_k , or it does).

Step 3: Write a bottom-up algorithm to compute the array values, following the recurrence.

Simply fill in the array values, following the recurrence.

$N[0, 0] := 0$

for $a := 1, 2, \dots, A$: $N[0, a] := \infty$

for $k := 1, 2, \dots, m$:

$N[k, 0] := 0$

for $a := 1, 2, \dots, A$:

$N[k, a] := N[k-1, a]$

if $c_k \leq a$ and $1 + N[k-1, a - c_k] < N[k, a]$:

$N[k, a] := 1 + N[k-1, a - c_k]$

Step 4: Use the computed values to reconstruct an optimal solution.

Idea: for every k, a , use coin c_k iff $N[k, a] \neq N[k-1, a]$.

if $N[m, A] = \infty$: return $\{\}$

$S = \{\}$

$a := A$

for $k := m, m-1, \dots, 1$:

if $N[k, a] \neq N[k-1, a]$:

$S := S \cup \{k\}$

$a := a - c_k$

return S

The worst-case runtime of the entire algorithm is $\Theta(mA)$, for filling in the values of $N[k, a]$ (reconstructing the solution takes only time $\Theta(m)$).