Worth: 5%

1. [20 marks]
   Consider the problem of multiplying two n-bit integers x, y, and suppose that n is a multiple of 3 (if not, we add one or two zeroes on the left to make this true). This time, suppose that we split each integer into three equal parts instead of two.

   (a) [5 points]
   State the exact relationship between x and each of its three parts X_2, X_1, X_0 (*i.e.*, give an equation that involves x and X_2, X_1, X_0).

   (b) [12 points]
   Write a divide-and-conquer algorithm to multiply two integers x, y based on this three-way split. Justify that your algorithm produces the correct answer, *i.e.*, show that the output of your algorithm is equal to $x \cdot y$.

   For full marks, your algorithm must run in time $o(n^2)$, *i.e.*, strictly less than $\mathcal{O}(n^2)$. Justify that this is the case by computing the running time of your algorithm (you may use the Master Theorem as long as you state clearly how it applies to your algorithm).

   (c) [3 points]
   Is your algorithm faster or slower than the divide-and-conquer algorithm shown in class with a running time of $\Theta(n^{\log_2 3})$?

2. [20 marks]
   Consider a variant on the problem of Interval Scheduling where instead of wanting to schedule as many jobs as we can on one processor, we now want to schedule ALL of the jobs on as few processors as possible.

   The input to the problem is (s_1, f_1), (s_2, f_2), $\cdots$, (s_n, f_n), where $n \geq 1$ and all $s_i < f_i$ are nonnegative integers. The integers s_i and f_i represent the start and finish times, respectively, of job i.

   A *schedule* specifies for each job i a positive integer $P(i)$ (the "processor number" for job i). It must be the case that if $i \neq j$ and $P(i) = P(j)$, then jobs i and j do not overlap. We wish to find a schedule that uses as few processors as possible, *i.e.*, such that $\max\{P(1), P(2), \cdots, P(n)\}$ is minimal.

   (a) [5 marks]
   Design an algorithm (write a pseudocode) to solve the above problem in time $o(n^2)$, *i.e.*, strictly less than $\mathcal{O}(n^2)$.

   (b) [10 marks]
   Prove that the above algorithm is guaranteed to compute a schedule that uses the minimum number of processors.

   (c) [5 marks]
   Briefly describe an efficient implementation of the algorithm, making it clear what data structures you are using. Express the running time of your implementation as a function of n (the number of jobs), using appropriate asymptotic notation.

3. [20 marks]
   Here is another variant on the problem of Interval Scheduling.

   Suppose we now have *two* processors, and we want to schedule as many jobs as we can. As before, the input is (s_1, f_1), (s_2, f_2), $\cdots$, (s_n, f_n) where $n \geq 1$ and all $s_i < f_i$ are nonnegative integers.

A *schedule* is now defined as a pair of sets (A_1, A_2), the intuition being that A_i is the set of jobs scheduled on processor i. A schedule must satisfy the obvious constraints: $A_1 \subseteq \{1, 2, \ldots, n\}$, $A_2 \subseteq \{1, 2, \ldots, n\}$, $A_1 \cap A_2 = \varnothing$, and for all $i \neq j$ such that $i, j \in A_1$ or $i, j \in A_2$, jobs i and j do not overlap.

   (a) [5 marks]
      Design an algorithm (write a pseudocode) to solve the above problem in time $o(n^2)$, *i.e.*, strictly less than $\mathcal{O}(n^2)$.

   (b) [10 marks]
      Prove that the above algorithm is guaranteed to compute an optimal schedule.

   (c) [5 marks]
      Briefly describe an efficient implementation of the algorithm, making it clear what data structures you are using. Express the running time of your implementation as a function of n (the number of jobs), using appropriate asymptotic notation.

4. [20 marks]
   Consider the problem of making change, given a finite number of coins with various values. Formally:

Input: A list of positive integer coin values c_1, c_2, $\cdots$, c_m (with repeated values allowed) and a positive integer amount A.

Output: A selection of coins $\{i_1, i_2, \ldots, i_k\} \subseteq \{1, 2, \ldots, m\}$ such that $c_{i_1} + c_{i_2} + \cdots + c_{i_k} = A$ and k is as small as possible. If it is impossible to make change for amount A exactly, then the output should be the empty set $\emptyset$.

   (a) [5 points]
      Describe a natural greedy strategy you could use to try and solve this problem, and show that your strategy does not work.

      (The point of this question is **not** to try and come up with a really clever greedy strategy — rather, we simply want you to show why the "obvious" strategy fails to work.)

   (b) [10 points]
      Give a detailed dynamic programming algorithm to solve this problem. Follow the steps outlined in class, and include a brief (but convincing) argument that your algorithm is correct.

   (c) [5 points]
      What is the worst-case running time of your algorithm? Justify briefly.

5. [20 marks]
   During the renovations at Union Station, the work crews excavating under Front Street found veins of pure gold ore running through the rock! They cannot dig up the entire area just to extract all the gold: in addition to the disruption, it would be too expensive. Instead, they have a special drill that they can use to carve a single path into the rock and extract all the gold found on that path. Each crew member gets to use the drill once and keep the gold extracted during their use. You have the good luck of having an uncle who is part of this crew. What's more, your uncle knows that you are studying computer science and has asked for your help, in exchange for a share of his gold!

   The drill works as follows: starting from any point on the surface, the drill processes a block of rock $10\text{cm} \times 10\text{cm} \times 10\text{cm}$, then moves on to another block 10cm below the surface and connected with the starting block either directly or by a face, edge, or corner, and so on, moving down by 10cm at each "step". The drill has two limitations: it has a maximum depth it can reach and an initial hardness

that gets used up as it works, depending on the hardness of the rock being processed; once the drill is all used up, it is done even if it has not reached its maximum depth.

The good news is that you have lots of information to help you choose a path for drilling: a detailed geological survey showing the hardness and estimated amount of gold for each 10cm × 10cm × 10cm block of rock in the area. To simplify the notation, in this homework, you will solve a two-dimensional version of the problem defined as follows.

Input: A positive integer d (the initial *drill hardness*) and two $[m \times n]$ matrices H, G containing non-negative integers. For all $i \in \{1,\ldots,m\}, j \in \{1,\ldots,n\}$, $H[i,j]$ is the *hardness* and $G[i,j]$ is the *gold content* of the block of rock at location i,j (with $i = 1$ corresponding to the surface and $i = m$ corresponding to the maximum depth of the drill).

There is one constraint on the values of each matrix: $H[i,j] = 0 \implies G[i,j] = 0$ (blocks with hardness 0 represent blocks that have been drilled already and contain no more gold).

Figure 1 below shows the general form of the input. Figure 2 shows a sample input.

Output: A drilling path $j_1, j_2, \ldots, j_\ell$ for some $\ell \le m$ such that:

- $\boxed{1 \le j_k \le n \text{ for } k = 1, 2, \ldots, \ell \text{ (\textbf{each coordinate on the path is valid});}}$
- $j_{k-1} - 1 \le j_k \le j_{k-1} + 1$ for $k = 2, \ldots, \ell$ (each block is underneath the one just above, either directly or diagonally);
- $H[1, j_1] + H[2, j_2] + \cdots + H[\ell, j_\ell] \le d$ (the total hardness of all the blocks on the path is no more than the initial drill hardness);
- $G[1, j_1] + G[2, j_2] + \cdots + G[\ell, j_\ell]$ is maximum (the path collects the maximum amount of gold possible).

Follow the dynamic programming paradigm given in class (the "five step" process) to give a detailed solution to this problem. Make sure to keep a clear distinction between each of the steps, and to explain what you are doing and why at each step — you will **not** get full marks if your answer contains only equations or algorithms with no explanation or justification.

$H[1,1],G[1,1]$	$H[1,2],G[1,2]$	$\cdots$	$H[1,n],G[1,n]$
$H[2,1],G[2,1]$	$H[2,2],G[2,2]$	$\cdots$	$H[2,n],G[2,n]$
$\vdots$	$\vdots$	$\ddots$	$\vdots$
$H[m,1],G[m,1]$	$H[m,2],G[m,2]$	$\cdots$	$H[m,n],G[m,n]$

Figure 1: General form of the input matrix.

2,2	**2,7**	0,0	2,3	1,8
2,2	**0,0**	1,2	2,0	1,1
1,4	1,1	**2,6**	2,1	3,8

Figure 2: Sample input with optimum path shown in **bold** (for $d = 4$).