

STA305/1004 - Review of Statistical Theory

January 8-9, 2020

Data

Experimental data describes the outcome of the experimental run. For example 10 successive runs in a chemical experiment produce the following data:

```
set.seed(100)
# Generate a random sample of 5 observations
# from a  $N(60, 10^2)$ 
dat <- round(rnorm(5, mean = 60, sd = 10), 1)
dat
```

```
## [1] 55.0 61.3 59.2 68.9 61.2
```

Distributions

Distributions can be displayed graphically or numerically.

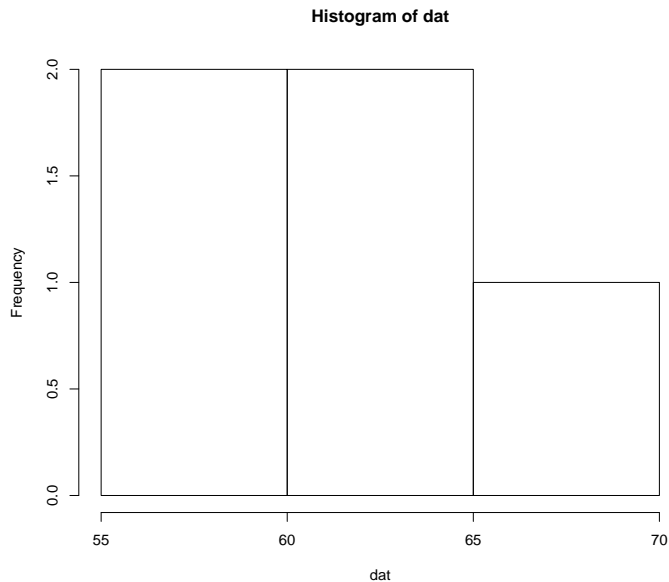
A histogram is a graphical summary of a data set.

```
summary(dat)
```

##	Min.	1st Qu.	Median	Mean	3rd Qu.	Max.
##	55.00	59.20	61.20	61.12	61.30	68.90

Distributions

```
hist(dat)
```



Distributions

- ▶ The total aggregate of observations that might occur as a result of repeatedly performing a particular operation is called a **population** of observations.
- ▶ The observations that actually occur are a **sample** from the population.

Continuous Distributions

- ▶ A continuous random variable X is fully characterized by its density function $f(x)$.
- ▶ $f(x) \geq 0$, f is piecewise continuous, and $\int_{-\infty}^{\infty} f(x)dx = 1$.
- ▶ The cumulative distribution function (CDF) of X is defined as:

$$F(x) = P(X \leq x) = \int_{-\infty}^x f(x)dx.$$

Continuous Distributions


- ▶ If f is continuous at x then $F'(x) = f(x)$ (fundamental theorem of calculus).
- ▶ The CDF can be used to calculate the probability that X falls in the interval (a, b) . This is the area under the density curve which can also be expressed in terms of the CDF:

$$P(a < X < b) = \int_a^b f(x)dx = F(b) - F(a).$$

- ▶ In R a list of all the common distributions can be obtained by the command `help("distributions")`.
- ▶ For example, the normal density and CDF are given by `dnorm()` and `pnorm()`.

Continuous Distributions

100 observations (using `rchisq()`) from a Chi-square distribution on 10 degrees of freedom χ^2_{10} . The density function of the χ^2_{10} is superimposed over the histogram of the sample.



Randomness

- ▶ A random drawing is where each member of the population has an equal chance of being selected.
- ▶ The hypothesis of random sampling may not apply to real data.
- ▶ For example, cold days are usually followed by cold days.
- ▶ So daily temperature not directly representable by random drawings.
- ▶ In many cases we can't rely on the random sampling property although design can make this assumption relevant.

Parameters and Statistics

What is the difference between a parameter and a statistic?

- ▶ A parameter is a population quantity and a statistic is a quantity based on a sample drawn from the population.

Example: The population of all adult (18+ years old) males in Toronto, Canada.

- ▶ Suppose that there are N adult males and the quantity of interest, y , is age.
- ▶ A sample of size n is drawn from this population.
- ▶ The population mean is $\mu = \sum_{i=1}^N y_i / N$.
- ▶ The sample mean is $\bar{y} = \sum_{i=1}^n y_i / n$.

Residuals and Degress of Freedom

$y_i - \bar{y}$ is called a residual.

- ▶ Since $\sum (y_i - \bar{y}) = 0$ any $n - 1$ completely determine the the last observation.
- ▶ This is a constraint on the the residuals.
- ▶ So n residuals have $n - 1$ degrees of freedom since the last residual cannot be freely chosen.

The Normal Distribution

The density function of the normal distribution with mean μ and standard deviation σ is:


$$\phi(x) = \frac{1}{\sigma\sqrt{2\pi}} \exp\left(\frac{-1}{2} \left(\frac{x - \mu}{\sigma}\right)^2\right)$$

The cumulative distribution function (CDF) of a $N(0, 1)$ distribution,

$$\Phi(x) = P(X < x) = \int_{-\infty}^x \phi(x) dx$$

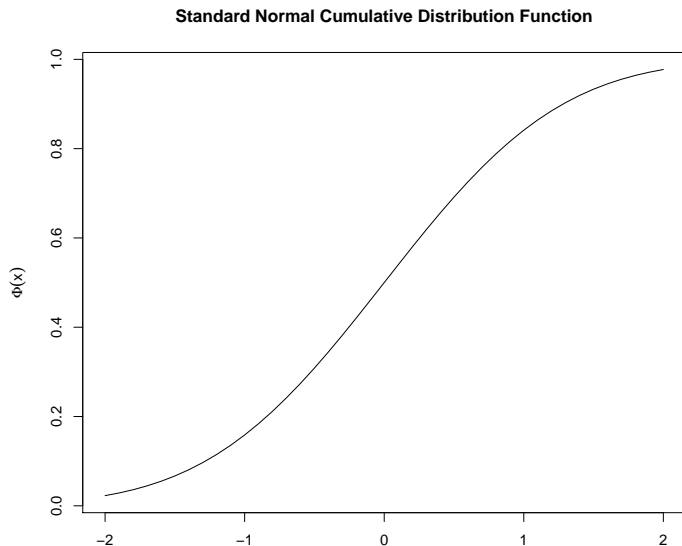
The Normal Distribution

```
x <- seq(-4,4,by=0.1)
plot(x,dnorm(x),type="l",main = "The Standard Normal Distribution",
     ylab=expression(paste(phi(x))))
```



The Normal Distribution

```
plot(x <- seq(-2,2,by=0.1),pnorm(x),type="l",  
     xlab="x",ylab=expression(paste(Phi(x))),  
     main = "Standard Normal Cumulative Distribution Function")
```



The Normal Distribution

A random variable X that follows a normal distribution with mean μ and variance σ^2 will be denoted by

$$X \sim N(\mu, \sigma^2).$$

If $Y \sim N(\mu, \sigma^2)$ then

$$Z \sim N(0, 1),$$

where

$$Z = \frac{Y - \mu}{\sigma}.$$

The Normal Distribution

$X \sim N(5, 3)$. Use R to find $P(4 < X < 6)$.

```
pnorm(6,mean = 5,sd = sqrt(3))-pnorm(4,mean = 5,sd = sqrt(3))
```

```
## [1] 0.4362971
```


Normal Quantile Plots

The following data are the weights from 11 tomato plants.


```
## [1] 29.9 11.4 26.6 23.7 25.3 28.5 14.2 17.9 16.5 21.1 24.3
```

Do the weights follow a Normal distribution?

Normal Quantile Plots

A normal quantile plot in R can be obtained using `qqnorm()` for the normal probability plot and `qqline()` to add the straight line.

```
qqnorm(tomato.data$pounds); qqline(tomato.data$pounds)
```



Central Limit Theorem

The central limit theorem states that if X_1, X_2, \dots is an independent sequence of identically distributed random variables with mean $\mu = E(X_i)$ and variance $\sigma^2 = \text{Var}(X_i)$ then

$$\lim_{n \rightarrow \infty} P\left(\frac{\bar{X} - \mu}{\frac{\sigma}{\sqrt{n}}} \leq x\right) = \Phi(x),$$


where $\bar{X} = \sum_{i=1}^n X_i/n$ and $\Phi(x)$ is the standard normal CDF. This means that the distribution of \bar{X} is approximately $N\left(\mu, \frac{\sigma}{\sqrt{n}}\right)$.

Central Limit Theorem


Example: A fair coin is flipped 50 times. What is the distribution of the average number of heads?

Central Limit Theorem

```
set.seed(100)  
Total.heads <- rbinom(100,50,0.5); Ave.heads <- Total.heads/50;  
hist(Ave.heads)
```



Central Limit Theorem



Chi-Square Distribution

Let X_1, X_2, \dots, X_n be independent and identically distributed random variables that have a $N(0, 1)$ distribution. The distribution of

$$\sum_{i=1}^n X_i^2,$$

has a chi-square distribution on n degrees of freedom or χ_n^2 .

The mean of a χ_n^2 is n with variance $2n$.

Chi-Square Distribution

Let X_1, X_2, \dots, X_n be independent with a $N(\mu, \sigma^2)$ distribution. What is the distribution of the sample variance $S^2 = \sum_{i=1}^n (X_i - \bar{X})^2 / (n - 1)$?

t Distribution

If $X \sim N(0, 1)$ and $W \sim \chi_n^2$ then the distribution of $\frac{X}{\sqrt{W/n}}$ has a t distribution on n degrees of freedom or $\frac{X}{\sqrt{W/n}} \sim t_n$.

t Distribution


Let X_1, X_2, \dots is an independent sequence of identically distributed random variables that have a $N(0, 1)$ distribution. What is the distribution of

$$\frac{\bar{X} - \mu}{\frac{S}{\sqrt{n}}}$$

where $S^2 = \sum_{i=1}^n (X_i - \bar{X})^2 / (n - 1)$?

t Distribution

Comparison of t Distributions



F Distribution


Let $X \sim \chi_m^2$ and $Y \sim \chi_n^2$ be independent. The distribution of

$$W = \frac{X/m}{Y/n} \sim F_{m,n},$$

where $F_{m,n}$ denotes the F distribution on m, n degrees of freedom. The F distribution is right skewed (see graph below). For $n > 2$, $E(W) = n/(n-2)$. It also follows that the square of a t_n random variable follows an $F_{1,n}$.

F Distribution

F Distributions



Linear Regression

Lea (1965) discussed the relationship between mean annual temperature and mortality index for a type of breast cancer in women taken from regions in Europe (example from Wu and Hammada).

The data is shown below.

```
#Breast Cancer data  
M <- c(102.5, 104.5, 100.4, 95.9, 87.0, 95.0, 88.6, 89.2,  
       78.9, 84.6, 81.7, 72.2, 65.1, 68.1, 67.3, 52.5)  
T <- c(51.3, 49.9, 50.0, 49.2, 48.5, 47.8, 47.3, 45.1,  
       46.3, 42.1, 44.2, 43.5, 42.3, 40.2, 31.8, 34.0)
```

Linear Regression

A linear regression model of mortality versus temperature is obtained by estimating the intercept and slope in the equation:

$$y_i = \beta_0 + \beta_1 x_i + \epsilon_i, i = 1, \dots, n$$

where $\epsilon_i \sim N(0, \sigma^2)$. The values of β_0, β_1 that minimize the sum of squares

$$\sum_{i=1}^n (y_i - (\beta_0 + \beta_1 x_i))^2,$$


are called the least squares estimators. They are given by:

- ▶ $\hat{\beta}_0 = \bar{y} - \hat{\beta}_1 \bar{x}$
- ▶ $\hat{\beta}_1 = r \frac{S_y}{S_x}$

r is the correlation between y and x , and S_x, S_y are the sample standard deviations of x and y respectively.

Linear Regression

```
plot(T,M,xlab="temperature",ylab="mortality index")
```




Linear Regression

```
reg1 <- lm(M~T)
summary(reg1) # Parameter estimates and ANOVA table
```

```
##
## Call:
## lm(formula = M ~ T)
##
## Residuals:
##      Min       1Q   Median       3Q      Max
## -12.8358  -5.6319   0.4904   4.3981  14.1200
##
## Coefficients:
##              Estimate Std. Error t value Pr(>|t|)
## (Intercept) -21.7947    15.6719  -1.391   0.186
## T             2.3577     0.3489   6.758 9.2e-06 ***
## ---
## Signif. codes:  0 '***' 0.001 '**' 0.01 '*' 0.05 '.' 0.1 ' ' 1
##
## Residual standard error: 7.545 on 14 degrees of freedom
## Multiple R-squared:  0.7654, Adjusted R-squared:  0.7486
## F-statistic: 45.67 on 1 and 14 DF,  p-value: 9.202e-06
```


Linear Regression

```
plot(T,M,xlab="temperature",ylab="mortality index")  
abline(reg1) # Add regression line to the plot
```



Linear Regression


```
#plot residuals vs. fitted  
plot(reg1$fitted,reg1$residuals);  
abline(h=0) # add horizontal line at 0
```



Linear Regression

#check normality of residuals

```
qqnorm(reg1$residuals); qqline(reg1$residuals)
```



Linear Regression

If there is more than one independent variable then the above model is called a multiple linear regression model.

$$y_i = \beta_0 + \beta_1 x_{i1} + \beta_2 x_{i2} + \cdots + \beta_k x_{ik} + \epsilon_i, \quad i = 1, \dots, n,$$

where $\epsilon_i \sim N(0, \sigma^2)$.

This can also be expressed in matrix notation as

$$y = X\beta + \epsilon$$

The least squares estimator is

$$\hat{\beta} = (X^T X)^{-1} X^T y.$$

The covariance matrix of $\hat{\beta}$ is $(X^T X)^{-1} \sigma^2$. An estimator of σ^2 is

$$\hat{\sigma}^2 = \frac{1}{n - k} \sum_{i=1}^n (y_i - \hat{y}_i)^2,$$

where $\hat{y}_i = \hat{\beta}_0 + \hat{\beta}_1 x_{i1} + \cdots + \hat{\beta}_k x_{ik}$ is the predicted value of y_i .

Weighing Problem

Harold Hotelling in 1949 wrote a paper on how to obtain more accurate weighings through experimental design.

Method 1

Weigh each apple separately.

Method 2

Obtain two weighings by

1. Weighing two apples in one pan.
2. Weighing one apple in one pan and the other apple in the other pan

Weighing Problem

Let w_1, w_2 be the weights of apples one and two. Each weighing has standard error σ . So the precision of the estimates from method 1 is σ .

If the objects are weighed together in one pan, resulting in measurement m_1 , then in opposite pans, resulting in measurement m_2 , we have two equations for the unknown weights w_1, w_2 :

$$w_1 + w_2 = m_1$$

$$w_1 - w_2 = m_2$$

Weighing Problem

This can also be viewed as a linear regression problem $y = X\beta + \epsilon$:

$$y = (m_1, m_2)', X = \begin{pmatrix} 1 & 1 \\ 1 & -1 \end{pmatrix}, \beta = (w_1, w_2)'.$$

Weighing Problem

The least-squares estimates can be found using R.

#step-by-step matrix multiplication example for weighing problem

```
X <- matrix(c(1,1,1,-1),nrow=2,ncol=2) #define X matrix
Y <- t(X)%*%X # multiply  $X^T$  by X ( $X^T * X$ ) NB: t(X) is transpose of X
W <- solve(Y) # calculate the inverse
W %*% t(X) # calculate  $(X^T * X)^{-1} * X^T$ 
```

```
##      [,1] [,2]
## [1,]  0.5  0.5
## [2,]  0.5 -0.5
```

*W # print $(X^T * X)^{-1}$ for SE*

```
##      [,1] [,2]
## [1,]  0.5  0.0
## [2,]  0.0  0.5
```